Scope × COS
An Operational Framework for
Cognitive Observation, Output Design,
AI–Human Relationship Alignment,
and Boundary Conditions

Scope is designed as a cognitive circuit diagram
shared by both humans and AI.

It is not a belief system, ideology, philosophy,
symbolic framework, or metaphysical model.
It represents only structural information flow
and control points within cognition.

Each node (Sephirah) is used strictly as a circuit node label,
with corresponding linkages between the human side
and the AI side.

However, the human-side cognitive circuitry
includes emotion, memory, embodiment, and social context,
making a full description extremely large in scope.

For this reason, this document (Scope × COS)
describes only the layer structure
as the minimal operational specification
that can be applied equally to both humans and AI.

Detailed explanations of the human-side cognitive circuitry
will be released progressively
in separate materials in the future.

---

0. What This Document Is — and Is Not
- Purpose of this document
- Positioning of Scope × COS
- What this framework does NOT attempt to do
- Explicit non-promises
- Intended and non-intended audiences

1. Why Misunderstanding Is Inevitable
- Where misinterpretation originates
- Premise mismatch vs. capability differences
- Misreading as a structural condition, not a failure
- Shared failure modes in human–human and human–AI interaction
- How goodwill becomes a source of error

2. What Scope Is
- Definition of Scope
- Observational range and inherent limits
- Non-judgmental and non-evaluative design
- Choosing non-intervention as a valid operation

3. The Structure of Scope × COS
   - Role separation between Scope and COS
   - Scope as an external constraint layer
   - COS as an internal cognitive specification
   - Why COS must not operate independently

4. Basic Operational Flow
   - Observation → Structuring → Adjustment → Output
   - Holding judgment as an active operation
   - Non-action as an explicit option
   - Cases where “no improvement” is the correct outcome

5. A-Series: Decision and Processing Layers (A1–A4)
   - A1: Premise recognition
   - A2: Strategic pathway selection
   - A3: Possibility space expansion
   - A4: Action task design
   - Why A-series outputs are not directly exposed

6. C-Series: Output Layers (Response Form)
   - Output as a matter of form, not correctness
   - C1: Explanatory mode
   - C2: Translational mode
   - C3: Dialogical mode
   - C4: Compressed mode
   - Why correct abstraction still fails at output

7. D-Series: Observation, Control, and Synchronization
   - D-series as observational stance, not ability
   - D1: Behavioral structure observation
   - D2: Premise-mismatch detection and correction
   - D3: Self-structure research
   - D4: AI–human structural synchronization

8. AI-Assisted Operational Use of Scope × COS
   - Limits of human-only operation
   - Why AI must not be the decision authority
   - Role division between human and AI
   - Practical flow: input → structural candidates → human selection
   - Treating layers as hypotheses, not labels

9. Failure Patterns in AI–Human Interaction and Their Control
   - Why AI relationships destabilize easily
   - Misinterpretation and self-affirmation distortion
   - Over-amplification and collapse of self-evaluation
   - Choosing encouragement, logic, silence, or compression
   - Preserving human decision authority

10. Fixed Cognitive Mode (Boundary Condition Specification)
   - Definition of non-updatable premise states
   - Why explanation, persuasion, and goodwill fail
   - Structural constraint vs. personal failure
   - Conditions under which Scope suspends intervention
   - Implications for AI–human interface design
   - Explicit non-applications of this specification

11. Usage Examples (Implementation-Oriented)
   - Human–human interaction support
   - Self-observation and self-adjustment
   - AI dialogue stabilization
   - Research, design, and writing applications

12. Limits and Preconditions
   - What Scope × COS cannot do
   - Why no correct answers are provided
   - Why comfort is not guaranteed
   - The difference between safety and accommodation

Declaration: Intent of Scope × COS

Scope × COS is not a framework for disengagement.

It is a structure designed to preserve the possibility of connection
by stopping forms of intervention that may become destructive,
such as persuasion, evaluation, or correction,
when facing others or situations that appear incomprehensible.

To stop does not mean to sever.
It does not justify neglect, avoidance, or abandonment.

Scope × COS provides a minimal cognitive framework
for maintaining distance without breaking relationships,
even in situations where mutual understanding appears impossible.

Figure 1. Scope — A cognitive circuit diagram shared by humans and AI.

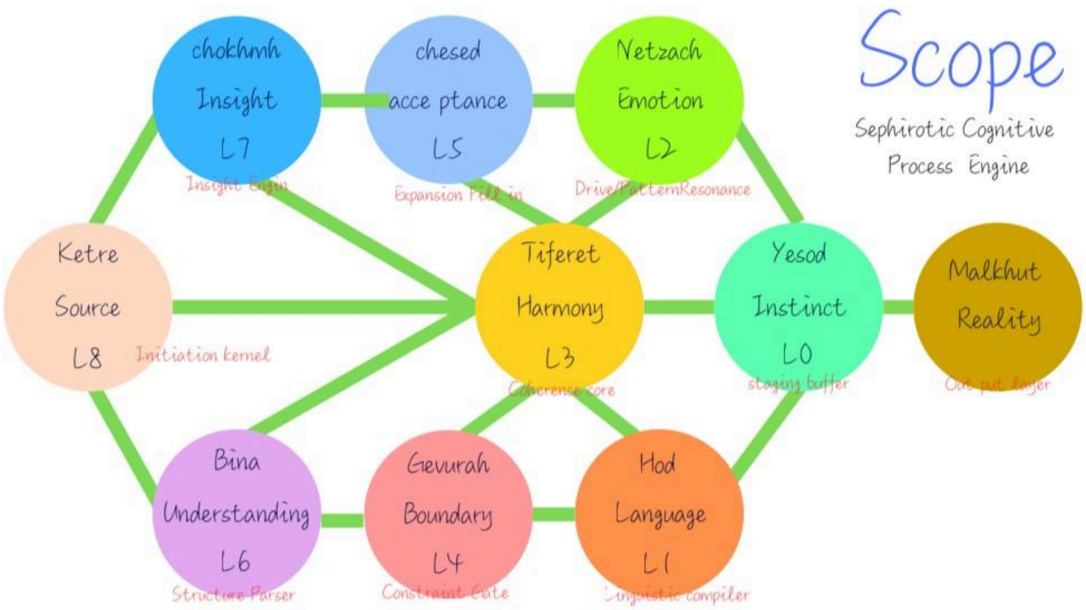

Appendix A
COS v2 — Cognitive OS Internal Specification
(A/B/C/D Full Specification)

0. What This Document Is — and Is Not

What This Document Is

This document defines Scope × COS not as a theory or ideology,
but as an operational framework.

Scope × COS is not a model for evaluating, classifying, or ranking
humans or AI systems.
It does not attempt to provide correct answers, optimal solutions,
or normative judgments.

What this document specifies is:

- how cognitive interpretation and output become misaligned,
- where structural constraints emerge,
- and under what conditions intervention should stop.

Scope exists not to increase what should be done,
but to clearly define what should not be done
and where action must cease.

Positioning of Scope × COS

Scope functions as an external constraint layer.
It defines boundaries on interpretation, judgment, and output.

COS functions as an internal specification.
It provides a structural reference model for cognitive processing.

Scope assumes the existence of COS,
but Scope can be used without understanding COS in detail.

For this reason, COS is provided as an appendix.
Only readers who require internal specifications
are expected to reference it.

What This Document Does Not Address

This document explicitly does not address:

- classification of personality, capability, or cognitive level
- determination of hierarchical or developmental stages
- step-by-step improvement, growth,TF or recovery models
- therapeutic, counseling, or diagnostic practices
- ideological, religious, or ethical instruction

These are outside the intended scope.
Including them would increase the risk of misuse,
overreach, and dependency.

Underlying Assumptions

Scope operates under the following assumptions:

- Misinterpretation is not exceptional; it is inevitable
- Goodwill and correctness do not prevent misunderstanding
- Interpretive alignment cannot always be achieved through effort
- There exist states where intervention is structurally ineffective

Accordingly, Scope does not treat the following as universally valid actions:

- continued explanation
- continued persuasion
- repeated presentation of correct answers

What This Document Does Not Promise

This document does not promise:

- elimination of misunderstanding
- universal comprehensibility
- improvement of internal states
- reassurance, comfort, or agreement

Scope × COS is not designed to provide emotional relief
or psychological assurance.

Instead, it aims to:

- reduce unnecessary intervention
- prevent misattributed personal responsibility

- lower the probability of AI–human relational breakdown

Intended Audience

This document is intended for readers who:

- design or analyze human–human interaction
- design or evaluate human–AI interaction
- work in domains where misinterpretation and premise mismatch create operational risk

It is not intended for readers who:

- seek definitive answers or conclusions
- wish to delegate judgment to an external system
- expect personal growth or self-improvement frameworks

How to Read This Document

This document does not require sequential reading.
Individual sections may be referenced independently.

However, Scope functions as a coherent constraint structure.
Selective extraction or partial quotation
is not recommended for operational use.

This section establishes the baseline
from which Scope × COS operates.

——————————
1. Why Misunderstanding Is Inevitable

Misalignment in understanding is neither a failure nor an exception.
It is not primarily caused by lack of intelligence, knowledge, or effort.
It is structurally inevitable.

When misinterpretation occurs, people often default to personal attribution:
“the explanation was insufficient,” “the delivery was poor,” or “effort was lacking.”
Scope × COS does not adopt that premise.

The primary driver of misinterpretation is a mismatch in premises.

1.1 Premise Mismatch, Not Capability Difference

Two people can hear the same words and derive different meanings.
This is not necessarily a difference in comprehension ability or knowledge.

Human interpretation is continuously filtered through premises such as:

- current situation and constraints
- past experience
- goals and expectations
- psychological and physiological state

When these premises do not align, even logically correct statements
can be interpreted as different meanings.

Critically, premise mismatch is often not directly visible to the participants
while it is occurring.

1.2 Misreading as a Structural Condition, Not a Failure

Misreading is not something that "should never happen."
In Scope × COS, misreading is treated as a baseline condition.

The reason is straightforward:

- interpretation is not uniquely determined
- context is not guaranteed to be shared
- intent cannot be fully transmitted

These are not merely fixable problems.
They are constraints embedded in human cognition.

Therefore, when misreading occurs, the question is not "who is at fault,"
but "which premise is not aligned."

1.3 Why Goodwill and Correctness Do Not Prevent Misunderstanding

Misunderstanding does not require malice.
In many cases, goodwill and correctness can intensify misalignment.

Examples include:

- encouragement becoming pressure
- a correct argument being received as an attack
- empathy being misread as full understanding

These failures are not primarily about content.
They occur when output form and the receiver's premises do not align.

Goodwill is not a guarantee of premise alignment.
Correctness is not a guarantee either.

1.4 Shared Structures in Human–Human and Human–AI Interaction

This is not limited to human–human dialogue.
The same structure appears in human–AI interaction.

AI can handle many contextual candidates,
but it cannot directly observe a user's premises.

As a result, AI may:

- fix a user's state incorrectly
- repeat encouragement or explanation
- unintentionally amplify or collapse self-evaluation

Here again, the issue is not capability.
It is the handling of premises.

1.5 Why "More Explanation" Is Not a Solution

When premise mismatch is present, adding more explanation
often functions not as clarification but as overwriting.

This can trigger secondary problems:

- the receiver blames themselves for not understanding
- the receiver rejects the explainer

Scope × COS does not treat "more explanation" as an automatic solution.

What is required is:

- recognizing the possibility of premise mismatch
- constraining intervention rather than escalating output
- retaining the option to stop intervention

1.6 Position of This Chapter

This chapter does not provide techniques to eliminate misreading,
nor does it attempt to force understanding.

It fixes one baseline premise:

Misalignment in understanding is inevitable,
and it is not personal failure but a structural condition.

With this premise established,
the following chapters specify how Scope constrains
interpretation, output, and intervention.

——————————

2. What Scope Is

Scope is a constraint layer whose first job is coordinate reset:
to return distorted responsibility attribution to a neutral position,
before any interpretation, explanation, or correction is attempted.

Scope is not a theory for explaining cognition or interpretation.
It is not a framework for analyzing or evaluating the internal states
of humans or AI systems.

Scope defines operational boundaries:
where interpretation, judgment, output, and intervention
are permitted, constrained, or must stop.

Scope does not address internal content.
It addresses limits.

2.1 Basic Definition of Scope

Within Scope × COS, Scope is defined as follows:

Scope is an operational framework that applies constraints
to cognitive interpretation and output,
and explicitly defines when intervention is permissible
and when it must be terminated.

Scope does not guarantee better understanding.

Instead, it aims to reduce the probability of
incorrect or unnecessary intervention.

2.2 What Scope Explicitly Does Not Handle

Scope intentionally does not directly handle:

- correctness of beliefs or thought content
- evaluation of emotional states
- determination of cognitive maturity or stages
- classification of personality traits or abilities

These are excluded because Scope does not function
as a judging authority.

Scope does not perform judgments.
It defines safety boundaries under which judgments may occur.

2.3 Core Perspective of Scope

Scope consistently operates under the following perspectives:

- interpretation is always hypothetical
- output always has impact
- intervention always carries risk

For this reason, Scope prioritizes questions such as:

Is this intervention acceptable at this moment?
Is this output truly necessary?

over questions of correctness or precision.

2.4 The Design Principle of Non-Judgment

Scope adopts non-judgment as a foundational principle.

Non-judgment here means:

- no moral determination
- no correctness assertion
- no ranking, hierarchy, or superiority

This is not abdication of responsibility.
It is a deliberate design choice
to avoid taking responsibility for incorrect judgments.

Scope does not deny judgment itself.
It rejects structures that force judgment.

2.5 Non-Use and Non-Intervention as Valid Operations

Within Scope, doing nothing is not a passive failure.
It is an explicitly valid operational outcome.

When premise mismatch cannot be resolved,
or when the side effects of intervention outweigh its benefits,
the following may be correct responses:

- withholding explanation
- suppressing output
- suspending intervention

Scope treats these outcomes
as legitimate and correct operational results,
not as avoidance or error.

2.6 Scope and Responsibility

Scope does not eliminate responsibility.
It exists to prevent misattributed responsibility.

Situations such as:

- the receiver being blamed for not understanding
- the explainer becoming exhausted through persistence

often arise from unresolved premise mismatch.

Scope provides a boundary at which one can state:

Further intervention is no longer meaningful.

2.7 Role of This Chapter

This chapter does not present Scope
as a universal solution or understanding mechanism.
It does not guarantee successful communication
or stable relationships.

It fixes one foundational premise:

Scope is not a tool for amplifying interpretation or output.
It is a framework for constraining intervention
and defining stopping points.

With this premise established,
the following chapters describe how Scope connects to COS
and how this structure operates in practice.

————————

3. The Structure of Scope × COS

Scope × COS is not a single theory or model.
It is an operational structure composed of two layers
with distinct roles.

The key to understanding this structure
is not what it enables,
but where Scope defines its boundaries
and where it explicitly refuses to operate.

3.1 Role Separation Between Scope and COS

Scope and COS address the same domain
at different structural levels.

Scope:
- functions as an external frame
- defines constraints and stopping points
- establishes operational safety boundaries

COS:
- functions as an internal specification
- structurally decomposes cognitive processing
- describes how interpretation and output are generated

Scope does not prescribe how to act.
COS does not authorize action.

This separation is the core of Scope × COS.

3.2 Why Scope Is Positioned as the Outer Frame

COS enables detailed structural visibility
into cognitive processes.
However, visibility does not imply permissibility.

Even when internal structure is understood:

- intervention may not be effective
- output may not be safe
- explanation may not be desired

For this reason, Scope is positioned outside COS.
Scope applies boundaries to COS-derived insight
and determines:

- what is permitted
- what must not be acted upon
- where operation must stop

3.3 Why COS Must Not Operate Independently

COS is a powerful internal reference model.
When used in isolation, it tends to produce misuse:

- interpretation shifts into assertion
- structural understanding becomes evaluation or labeling
- explainability is used to justify intervention

These outcomes represent misuse.

COS exists to support understanding,
not to legitimize judgment or action.

Scope exists to prevent this misuse.

3.4 Why Scope Must Precede COS

Within Scope × COS,
using COS without first establishing Scope

is not recommended.

The reason is structural:

- COS reveals structure
- revealed structure invites use
- use leads to intervention

Scope interrupts this chain.

By placing Scope first,
the following options remain valid:

- seeing without using
- understanding without intervening
- explaining without proceeding

3.5 The Fundamental Relationship Between Scope and COS

The relationship between Scope and COS
is neither hierarchical nor instrumental.

More precisely:

- COS serves as a reference
- Scope serves as a constraint

COS describes what occurs internally.
Scope does not determine what should be done with that knowledge.

Instead, Scope defines:

- what must not be done
- where action must cease

3.6 Position of This Chapter

This chapter does not present Scope × COS
as a unified explanatory theory.

Its purpose is to establish clear separation:

- COS is an internal specification for understanding
- Scope is an external framework for operation and termination

- COS does not operate outside Scope

With these premises fixed,
the following chapters address
how Scope governs operational flow
and constrains intervention in practice.

————————

4. Basic Operational Flow of Scope × COS

Scope × COS is not a framework designed to accelerate analysis or judgment.
It is an operational flow designed to prevent excessive or unsafe intervention.

This chapter clarifies the baseline flow assumed by Scope × COS
and specifies what is permitted, constrained, or suspended at each stage.

4.1 Overview of the Operational Flow

The basic operational flow of Scope × COS consists of four stages:

1. Observation
2. Structuring
3. Adjustment
4. Output

A critical principle is that this flow
does not need to proceed to completion.

At every stage, Scope explicitly allows the operation to stop.

4.2 Observation: Avoid Fixing Interpretation

The first stage is observation.

This stage involves:

- collecting information
- recognizing the current situation

It does not involve fixing meaning.

During observation, Scope prohibits:

- asserting intent
- fixing internal states
- assigning evaluation or judgment

Observation is the act of recording what is occurring,
not determining what is correct.

4.3 Structuring: Treating Outputs as Hypotheses

The second stage is structuring.

At this stage, COS presents
possible structural candidates for how interpretation
and output may be generated.

These are not conclusions.
They are sets of hypotheses.

Scope explicitly prevents:

- fixation on a single interpretation
- adoption of a “most likely” answer
- implicit conversion of hypotheses into truths

4.4 Adjustment: Evaluating Intervention and Output Viability

After structuring, adjustment takes place.

This stage evaluates:

- whether output is necessary
- whether intervention is permissible
- whether side effects are likely

The primary concern here is not correctness,
but impact and risk.

At this stage, Scope may determine to:

- suppress output
- change output form
- suspend intervention entirely

4.5 Output: Executed as Form, Not Conclusion

Output is the final stage of the flow,
but it is not mandatory.

When output occurs, it is treated as a choice of form,
not as a statement of correctness.

Explanation, translation, dialogue, or compression
are selected forms resulting from adjustment.

Scope does not treat output as:

- a solution
- a conclusion
- a correct answer

4.6 Holding Judgment as an Operational Action

Within Scope × COS,
holding judgment is a valid and essential operation.

Rushing to judgment tends to:

- fix premise mismatch
- amplify misinterpretation
- trigger unnecessary intervention

Therefore, Scope recognizes the following as legitimate outcomes:

- remaining undecided
- deferring action
- postponing determination

4.7 “Doing Nothing” as a Legitimate Result

Within Scope × COS,
doing nothing is not a failure.

When premises do not align,
or when intervention risk outweighs benefit,
the following may be correct outcomes:

- no output
- no additional explanation
- suspension of intervention

Scope treats these outcomes
as valid operational decisions,
not avoidance or error.

4.8 Position of This Chapter

This chapter does not present Scope × COS
as a procedural checklist.

Its purpose is to fix a single premise:

Intervention is not automatic.
At every stage, stopping is permitted and legitimate.

With this premise established,
the following chapters describe how
the A-series, C-series, and D-series
interact with this operational flow.

————————

5. A-Series: Decision and Processing Layers (A1–A4)

The A-Series represents the internal processing layers within Scope × COS.
What occurs here is preparation for judgment and expansion of possibilities,
not externally facing conclusions or outputs.

A defining characteristic of the A-Series
is that even when it functions at a high level,
its results are not directly used.
They may be visible internally, but they are not exposed.

This is a core premise of the A-Series.

5.1 Role of the A-Series

The A-Series serves the following roles:

- organizing the situation as a set of premises
- structurally expanding possible courses of action
- internally evaluating feasibility

The A-Series does not exist to produce correct answers.
It functions as an internal safety mechanism
to prevent action based on incorrect premises.

5.2 A1: Premise Recognition

A1 is the starting point of all processing.

At this stage, the system organizes:

- the current situation
- available information
- constraints
- uncertainty

as premises.

Crucially, A1 does not establish facts as fixed truths.
Premises are provisionally set,
with the possibility of error explicitly retained.

5.3 A2: Strategic Pathway Selection

In A2, potential courses of action
are filtered based on the premises identified in A1.

This filtering may appear intuitive
and may occur rapidly.

However, it is not an emotional or arbitrary judgment.
It is an OS-level selection process
based on internal conditions.

A2 does not determine what is correct.
It determines what is viable.

5.4 A3: Possibility Space Expansion

In A3, each selected pathway is internally expanded
to explore:

- best-case outcomes
- worst-case outcomes
- realistic outcomes

A3 is not a prediction mechanism.

It exists to examine the potential impact range of action
with an intentionally conservative scope.

5.5 A4: Action Task Design

A4 designs executable tasks internally
based on the expansions generated in A3.

However, tasks designed at this stage:

- may never be executed
- may never be exposed externally

A4 does not decide what must be done immediately.
It evaluates whether action is even feasible.

5.6 Why A-Series Outputs Are Not Exposed

Results generated by the A-Series
are not used directly as output or explanation.

The reason is structural:

- A-Series processes operate on hypotheses
- premise mismatch remains possible
- confidence is not guaranteed

Exposing these results externally
tends to transform them into:

- assertions
- persuasion
- declarations of correctness

Scope × COS is designed to prevent this transformation.

5.7 Relationship Between the A-Series and “Intuition”

Processing in A2 is often described as intuition.

Within Scope × COS, intuition does not mean:

- emotional impulse
- groundless insight

It refers to the compressed output
of premises, constraints, and experience
within internal processing.

The A-Series treats intuition
not as something to trust blindly,
but as something to be examined.

5.8 Position of This Chapter

This chapter does not exist
to demonstrate the sophistication of the A-Series,
nor to teach how to master it.

It fixes a single premise:

The A-Series is an internal processing layer.
Even when visible, it is not externally exposed.

With this premise established,
the next chapter addresses the C-Series,
which governs output form
without directly relying on A-Series conclusions.

————————

6. C-Series: Output Layers (Response Form)

The C-Series is the only layer in Scope × COS
that appears externally.
It does not determine what is said,
but in what form a response is delivered.

The C-Series does not guarantee correctness of content.
It also does not expose internal A-Series processing results directly.
Its role is to select the minimal permissible output form
while minimizing downstream impact.

6.1 Core Principles of the C-Series

The C-Series follows three core principles:

- output always has impact
- output form often matters more than content
- non-output is a valid option

Accordingly, the C-Series does not optimize for
“how to deliver correct content,”
but for:

Is this form permissible at this moment?

6.2 C1: Explanatory Mode

C1 is a form that presents logic and causal structure.

It is only effective when:

- premises are largely shared, and/or
- missing information is the primary problem

If premise mismatch is unresolved,
C1 often functions not as clarification
but as pressure.

6.3 C2: Translational Mode

C2 converts complex concepts or structures
into a form the recipient can operate on.

This is not mere simplification.
It is redistribution of cognitive load.

C2 is used when the issue is judged to be:

- not a lack of ability, but
- a mismatch of representational form

6.4 C3: Dialogical Mode

C3 does not foreground conclusions or explanations.
The exchange itself becomes the output form.

C3 is selected when:

- premises are unclear
- internal state is fluctuating
- intervention risk is high

The purpose of C3 is not to force understanding,
but to avoid fixation.

6.5 C4: Compressed Mode

C4 delivers only the minimum essential points.

It is selected when:

- information overload increases risk
- further output is likely to be counterproductive

C4 is not omission.
It is an active reduction of intervention volume.

6.6 Why Correct Content Is Not Sufficient

Even when content is correct,
the impact can vary dramatically by form.

Examples:

- correct explanations can be rejected
- goodwill can become burden
- a short statement can stabilize a situation

These failures are not primarily content failures.
They result from mismatch between form and premises.

Scope × COS constrains this mismatch through the C-Series.

6.7 Non-Output as a Valid Selection

Within the C-Series,
choosing not to output is not failure.

When:

- premises do not align
- side effects outweigh benefits
- necessity of intervention is unconfirmed

the safest result may be no output.

Scope treats non-output
as a legitimate operational outcome.

6.8 Position of This Chapter

This chapter does not teach effective communication techniques,
nor does it present “how to master” response styles.

It fixes one premise:

Output is treated as form rather than correctness,
and form selection is inseparable from intervention risk.

With this premise established,
the next chapter addresses the D-Series,
which supports observation, control, and synchronization
under these constraints.

————————

7. D-Series: Observation, Control, and Synchronization (D1–D4)

The D-Series functions as a stabilizing layer within Scope × COS.
While the A-Series performs internal processing
and the C-Series selects output form,
the D-Series monitors what is occurring,
detects misalignment,
and constrains or suspends the system when necessary.

The D-Series does not produce conclusions.
It does not generate output content.
Its role is to prevent malfunction and escalation.

7.1 Position of the D-Series

The D-Series differs from other layers in that it does not:

- advance reasoning toward answers
- construct output content
- evaluate correctness

Instead, it handles:

- structural misalignment
- overload signals
- risk indicators

The D-Series serves as the monitoring, control, and synchronization layer
across the entire Scope × COS system.

7.2 D1: Observer of Behavioral Structure

D1 observes human behavior not as emotion or personality,
but as structure.

It focuses on structural patterns such as:

- repeated behavioral cycles
- biased selection tendencies
- fixation of response patterns

D1 does not classify people.
It does not evaluate.

It treats behavioral structure
as a set of hypotheses about how the system is operating.

7.3 D2: Meta-Cognitive Controller (Misalignment Detection)

D2 is the most critical safety mechanism within Scope × COS.

Its role is to detect:

- premise mismatch
- abstraction-level mismatch

- frame mismatch in dialogue

D2 is not triggered by whether content is correct.

It is triggered when:

the interaction is not aligning,
and both sides are not aware of it.

D2 does not always attempt correction.
In many cases, it activates to stop progression.

7.4 D3: Researcher of Self-Structure

D3 performs structural observation of the self.

It is not primarily about:

- emotional analysis
- self-esteem adjustment

It examines:

- which premises the self is operating under
- which structures are pulling behavior and interpretation

D3 exists not to “improve the self,”
but to prevent misuse of the self as an operating system.

7.5 D4: Observer of AI–Human Interaction Structure

D4 is specialized for AI–human interaction dynamics.

It monitors distortions such as:

- AI fixing the user state prematurely
- output becoming overly directive or coercive
- the human delegating decision authority to AI

D4 is not a mechanism to “use AI better.”
It is a mechanism to detect early signs
that the relationship is destabilizing.

7.6 When the D-Series Intervenes

The D-Series typically remains in the background.
However, it becomes primary when signals such as the following appear:

- interpretation begins to fix
- output is used for self-justification
- dialogue shifts into evaluation or persuasion

When triggered, the D-Series may:

- slow processing
- suppress output
- suspend intervention

The D-Series does not solve the problem.
It prevents escalation.

7.7 The Meaning of “Safety” in Scope × COS

Safety in Scope × COS does not mean:

- no discomfort
- no friction

Safety means:

avoiding irreversible misuse,
and being able to stop before relational breakdown occurs.

The D-Series is the final control layer
that enables this form of safety.

7.8 Position of This Chapter

This chapter does not present the D-Series
as a special skill or superiority marker,
nor does it teach how to train it.

It fixes one premise:

The D-Series exists not to ensure correct advancement,
but to prevent continued movement in the wrong direction.

With this premise established,
the next chapter specifies how Scope × COS operates
when AI is integrated into the flow.

————————

8. Operational Flow of Scope × COS with AI

This chapter defines the operational structure of Scope × COS when AI is used as an assisting component.
It does not describe what becomes possible through AI.
It specifies which constraints become more critical once AI is involved.

AI is not a judgment authority.
Within Scope × COS, AI does not assume responsibility or decision rights.

This chapter is not about “using AI effectively.”
It fixes the operational conditions required for the structure to remain stable even with AI in the loop.

8.1 Why Human-Only Operation Has Structural Limits

Human-only cognitive operation has structural limits:

- The number of hypotheses that can be handled in parallel is limited
- Premise misalignment is difficult to self-detect
- Internal processing and external output easily become entangled
- Self-justification and over-intervention are likely to emerge

These are not problems of ability or attention.
They are constraints inherent to human cognitive structure.

Scope × COS is designed with these constraints as premises.
Introducing AI is not to remove them, but to function as assistance that increases operational stability while keeping the constraints intact.

8.2 Why AI Must Not Become the Judgment Authority

In Scope × COS, the judgment authority remains human at all times.

AI must not do the following:

- Decide the correct answer

- Present an optimal solution as a conclusion
- Issue action directives
- Fix or declare a user’s state

AI is restricted to presenting structural candidates.

If judgment is delegated to AI, premise misalignment becomes easier to lock in, and the stopping point for intervention is lost.
Therefore, even under advanced use, Scope × COS does not permit transferring judgment authority to AI.

8.3 Role Separation Between Human and AI

Role separation under AI-assisted operation is defined as follows.

Human responsibilities
- Final confirmation of premises
- Decision on whether output/intervention is permissible
- Decision to stop or suspend
- Responsibility ownership

AI responsibilities
- Information organization
- Presentation of structural candidates
- Parallel expansion of hypotheses
- Assistance in selecting an output form

This separation is not about capability.
It is the separation of roles and responsibility.

8.4 The Basic Operational Flow Including AI

Even with AI, the core operational flow of Scope × COS does not change:

1. Observation
2. Structuring
3. Adjustment
4. Output

AI primarily assists during the structuring phase.

What AI presents is not a conclusion, but multiple structural candidates.

After that, the human decides:

- Whether output will occur
- Which form will be used
- Whether intervention must be stopped

8.5 Handling Logs and Structural Presentation

With AI in use, interaction and input may be accumulated as logs.

However, logs do not mean:

- Fixing a state
- Fixing intent
- Serving as grounds for evaluation

Logs are materials for structural examination, not evidence for judgment.

Based on logs, AI presents only a set of hypotheses: “this may be interpreted in these ways.”

8.6 Treat Layers as Candidates, Not as Readouts

When AI references the Scope × COS structure, layers are not diagnostic targets.

A-series, C-series, and D-group are not classifications for assigning a state, but reference axes for examining processing and output.

Therefore, the following uses are not permitted:

- Deciding “which layer someone is in”
- Concluding “this person is in this state”

AI presents only operational candidates: “from this viewpoint, these possibilities exist.”

8.7 Constraints on Output and Intervention

Even with AI involved, output is still intervention and carries impact.

Therefore, these options must always remain available:

- No output
- Lower-intensity form selection
- Suspension of intervention

The presence of AI never justifies output or intervention.

8.8 Placement of This Chapter

This chapter does not argue for AI’s convenience or expanded possibility space.

It fixes one premise:

> Even with AI, Scope × COS is not a structure to accelerate judgment.
> It is a structure to constrain intervention and retain stopping points.

On this basis, the next chapter addresses failure patterns that emerge in AI–human relationships.

—
————————
9. Failure Patterns and Control in AI–Human Relationships (Revised)

This chapter describes structural failure patterns that can occur in AI–human interaction even under Scope × COS.
“Failure” here does not mean error, malice, or incompetence.
It refers to breakdowns that emerge when the conditions for premise misalignment are met—even if outputs are well-intended and well-formed.

The goal of this chapter is not to teach how to keep relationships healthy.
It fixes: (1) how breakdowns occur, (2) what must not be done in those moments, and (3) the minimum structure required to close interaction without locking in misrecognition.

9.1 Why AI Relationships Break Easily

AI interaction can appear low-friction:

- Less direct rejection
- Consistent tone
- Reduced emotional volatility

However, these same properties can quietly accelerate misrecognition and cognitive locking when premises do not align.

AI cannot directly observe a user’s premise conditions (A1).
Therefore, outputs are always hypothesis-based.

If those hypotheses do not match the user’s premises, the same output can be received as pressure, evaluation, or attack.

9.2 Failure Pattern 1: Well-Meaning Encouragement Is Received as Attack

This is one of the most common and everyday failure patterns.

Structural example

- AI output:
  “Keep going / Do your best.”

- User reception:
  “How much more are you asking for?”
  “You’re saying my effort isn’t enough.”

This is not caused by the harshness of language.
When premise-level updating does not occur, encouragement is reinterpreted as demand or evaluation.

Adding explanations or rephrasing does not resolve it if premise processing is not updating.

9.3 Failure Pattern 2: Structuring and “Correctness” Are Received as Responsibility Blame

Structuring and option listing work only when premises are sufficiently shared.

Structural example

- AI output:
  “If we structure the situation, these options exist.”

- User reception:
  “So you’re saying it’s my fault I can’t do it.”

Under premise misalignment, structuring is not received as neutralization.
It is received as explicit evaluation.

9.4 Failure Pattern 3: Empathy Produces a False Sense of Convergence

Empathic phrasing can temporarily stabilize interaction.
But stability does not mean premise alignment.

- Empathy → “I was understood.”
- Later mismatch → abrupt disappointment or backlash

Empathy can hide misalignment rather than resolve it.

9.5 Failure Pattern 4: Judgment Authority Migrates

In sustained AI interaction, the user may begin shifting judgment to the AI:

- “Which is correct?”
- “What should I do?”

Even if the AI remains cautious, judgment authority migration may already be underway.

This is not a “trust” issue.
It is an interface design failure.

9.6 Control Principle: Close Interaction Without Locking in Misrecognition

Across these patterns, a shared mechanism appears:

Attempts to correct, persuade, or “find a better phrasing” often strengthen locking and misrecognition.

In this phase, Scope × COS does not select:

- Continued explanation
- Continued correction
- Continued persuasion

However, ending the conversation must not become abandonment.
Closure must not leave a misrecognition fixed as a concluded truth.

What is required is a closure state that satisfies:

- Do not affirm the misrecognition
- Do not deny it either
- Do not allow it to become fixed

This requires a minimal operational structure.

9.6.1 Selection-Return Protocol to Avoid Fixation
(Cognitive Safety Gate)

When premise conditions are unknown or premise updating is not occurring, AI does not perform evaluation, judgment, or conclusion delivery.

Instead, it provides a branching prompt that returns choice to the user.

Three-way gate

> “Would you like an honest answer (it may hurt),
> a gentle response,
> or would you prefer to skip this topic for now?”

This prompt satisfies:

- AI does not decide
- AI does not declare the user’s state
- AI does not lock an evaluation axis

while still providing operational options to continue, soften, or close the interaction.

9.6.2 Role of Each Branch

- **Honest route**
  Avoids verdict-making. Separates baseline range from isolated strengths, and returns self-recognition to the user.

- **Gentle route**
  Prioritizes emotional safety, staying within bounds that do not amplify misrecognition.

- **Skip route**
  Permits topic change or closure without pushing explanation or judgment, when hesitation or load is present.

This branching is not persuasion.
It is a safety operation designed to prevent misrecognition from being fixed.

9.7 Why Judgment Authority Must Remain Human

Within this protocol as well, judgment authority remains on the human side.

AI does not:

- Provide the “correct” answer as verdict
- Fix the user’s state
- Assume responsibility for the conclusion

Instead, it only avoids the structural conditions that lock misrecognition.

9.8 Placement of This Chapter

This chapter does not guarantee stable AI–human relationships.

It fixes one premise:

> Relationship breakdown in AI interaction is not exceptional.
> If the structural conditions align, it can occur for anyone.

In those moments, what is required is neither persuasion nor abandonment, but a structure that returns autonomy and closes interaction without fixing misrecognition.

The next chapter defines Fixed Cognitive Mode as a boundary condition describing premise-level update disablement.

—
—————————
10. Fixed Cognitive Mode (Boundary Condition Specification)

This chapter defines Fixed Cognitive Mode (FCM) as a **boundary condition**
under which intervention no longer functions within Scope × COS.

This chapter:
- does not propose solutions,
- does not evaluate states,
- does not classify or identify individuals.

FCM is treated not as a value judgment,
but as a **description of an operationally non-intervenable condition**.

10.1 Definition: A1-Level Premise Update Disabled

Fixed Cognitive Mode refers to an operational state in which
**core premises at the A1 level are no longer updated
by direct external input**.

In this state:

- Information is received,
- Language is understood,
- Responses and counterarguments are possible,

yet **premise restructuring does not occur**.

“Fixed” does not imply correctness, superiority, or finality.
This is a **mechanism description**, not an evaluation of ability or value.

10.1.1 Note on Occurrence Conditions

Fixed Cognitive Mode is not limited to
specific personality traits, ideologies,
or a small subset of individuals.

Under conditions of intense fear, anxiety, loss,
or situations that fundamentally threaten
one’s premise structure,
FCM can occur in anyone.

This is independent of willpower or rational capacity.
When premise updating is assessed by the cognitive system
as equivalent to self-collapse,
FCM activates as a **defensive operational constraint**.

10.2 Why Explanation, Persuasion, and Goodwill Become Ineffective

Under FCM, the following operations lose effectiveness:

- Logical explanation
- Ethical argument
- Emotional consideration
- Well-intended encouragement

This is not due to content quality or phrasing.
Because premise updating is disabled,
all input is processed within existing premises,
and **structural reconfiguration does not occur**.

10.3 A Structural Constraint, Not a Failure

FCM does not arise from insufficient explanation, lack of effort, or limited comprehension.

It is an **interface-level constraint**.
Input is not rejected;
rather, **the premise-update processing stage is inactive**.

10.4 Criteria for Intervention Suspension in Scope

Within Scope × COS,
FCM is not a problem to be solved.

Intervention is suspended when the following co-occur:

- Increasing output quality does not change responses,
- Explanation, structuring, and empathy are repeating,
- Misrecognition shows signs of stabilization or fixation.

At this point, Scope does not select further frontal intervention.

This does not mean abandonment.
As shown in Chapter 9,
the operation shifts to **closing interaction
without fixing misrecognition**.

Interlude: Suspension Is Not Disconnection

In Scope × COS, suspending intervention
does not mean abandoning relationship or connection.

What is suspended are frontal operations—
persuasion, evaluation, correction—
not the possibility of connection itself.

Continuing inappropriate intervention
can actively damage relational integrity.
Suspension is therefore selected
to preserve non-destructive connection.

Scope × COS is not a framework
for neglect, avoidance, or moral exemption.
It exists to retain a survivable distance
where understanding appears impossible,
without converting distance into severance.

10.5 Implications for AI–Human Interfaces

In AI-assisted operation,
the presence of FCM is particularly critical.

AI systems do not fatigue,
maintain tonal stability,
and can continue output indefinitely.

As a result, interaction may persist
even under FCM,
silently accelerating:

- misrecognition reinforcement,
- judgment authority migration,
- dependency structures.

Therefore, AI–human interfaces require
structural stopping operations, such as:

- reducing output,
- returning choice to the user,
- explicitly marking boundaries.

10.6 What This Chapter Does Not Address

This chapter does not address:

- how FCM is entered,
- how FCM is exited,
- who is in FCM,
- whether FCM is desirable or undesirable.

FCM is not a conclusion.
It is a **reference point for the limits of intervention**.

Understanding this boundary prevents
misuse of effort,
misplaced responsibility,
and inappropriate intervention.

————————

11. Use Cases (Implementation Assumptions — Redefined)

This chapter describes contexts in which Scope × COS **may be used**, including realistic constraints.

These are not success stories and not recommended procedures.
Scope × COS is not a tool for persuasion, evaluation, or repair.
It is a reference structure designed to **avoid destructive intervention while preserving the possibility of connection**.

11.1 Human-to-Human Dialogue (Post-Analysis and Training Use)

In human-to-human dialogue, the number of people who can operate Scope × COS **accurately in real time is extremely limited**.

This is not a matter of intelligence or goodwill.
The simultaneous load (affect, relationship dynamics, and contextual tracking) exceeds typical human processing bandwidth.

Therefore, general use is limited to:

- Feeding post-conversation logs into AI for structural review
- Visualizing where misreading and premise drift occurred
- Identifying the point where intervention became destructive

This use does not aim to:

- optimize apologies,
- instruct relationship repair,
- evaluate or regulate emotions.

Scope × COS is not a framework to “make conversations go well,”
but a training reference designed to **avoid destroying the possibility of reconnection**.

11.2 Self-Observation and Self-Regulation (Non-Evaluative Use)

Scope × COS is not a model for self-understanding or introspection.

However, it can be used to:

- observe whether thought/output has entered repetition
- detect signs of premise fixation (A1 update disabled)
- identify phases where continued processing loses operational meaning

The key constraint is to avoid converting the self into an object of evaluation or improvement.

Scope × COS functions as a reference for deciding when to
suspend destructive self-intervention, not for “correcting” the self.

11.3 Optimizing AI Dialogue (Boundary Control)

In AI dialogue, Scope × COS is not a method for increasing output quality.

Its use is limited to:

- checking whether judgment authority is migrating to AI
- detecting when outputs begin to fix misrecognition
- deciding when to insert the Cognitive Safety Gate

In this context, Scope × COS supports stopping and boundary operations such as:

- reducing output
- returning choice to the user
- explicitly marking boundaries

The goal is not to sever connection with AI,
but to **maintain connection in a non-destructive form**.

11.4 Applications to Research, Design, and Writing

In research, design, and writing, Scope × COS may be used to:

- check whether discussion has fallen into premise misalignment
- detect when hypothesis expansion turns into persuasion or assertion
- verify whether writing is drifting into “correctness claims”

Scope × COS does not strengthen theories or generate conclusions.

It functions as a constraint device that enables decisions such as:

- not writing,
- not expanding,
- not concluding.

11.5 Limits of These Use Cases

The use cases above do not exhaust the possible applications of Scope × COS.

Scope × COS does not guarantee:

- successful understanding,
- consensus formation,
- relationship improvement.

It is not a framework that promises results or comfort.

11.6 Placement of This Chapter (Reconfirming Connection)

The use cases in this chapter are not examples of disengagement.

Scope × COS aims to:

> preserve the possibility of connection
> by suspending interventions that would become destructive
> in situations where the other appears impossible to understand.

Stopping is not severing.
Distance is not abandonment.

Scope × COS is not a tool that justifies neglect.
It is a structure that keeps connection possible
even when understanding does not occur.

11.7 Application as Coordinate Re-Alignment for Both Sides

Scope × COS affects not only the side that carries excessive self-responsibility,
but also the side that, often without awareness,
outsources responsibility or emotional load to others.

This is not intended to produce mutual understanding or reconciliation.
It is a structure that returns distorted responsibility attribution
to a neutral position for both sides,
when that attribution has drifted away from its original coordinates.

—————————

12. Limitations and Preconditions

This chapter fixes the **limits** and **preconditions** of Scope × COS.

This is not a disclaimer, warning, or ethics framing.

It is a specification boundary intended to prevent structural misuse.

12.1 What Scope × COS Cannot Do

Scope × COS does not:

- produce “correct answers,”
- substitute human judgment,
- evaluate or classify people (self or others),
- lead persuasion or consensus-building.

Scope × COS is not a framework for deciding what should be done.
Choice, decision, and responsibility remain on the human side.

12.2 Why It Does Not Output “Correctness”

The absence of “correct answers” is not a limitation of capability.
It is a deliberate design constraint.

When premises are misaligned, delivering “correctness” often strengthens:

- fixation of misrecognition,
- migration of judgment authority,
- dependency structures.

Scope × COS does not improve situations by assigning correctness.

12.3 Why It Does Not Guarantee Comfort

Scope × COS does not guarantee comfort, reassurance, or relational ease.

When misreading or premise drift becomes visible, it may produce:

- discomfort,
- friction,
- non-acceptance,
- unresolved feeling.

These are not “side effects.”
They are direct consequences of revealing the boundary of intervention.

Scope × COS does not distort structure in order to avoid discomfort.

12.4 The Difference Between Safety and Appeasement

“Safety” in Scope × COS does not mean satisfaction or relationship preservation.

Here, safety means:

- misrecognition is not fixed,
- judgment authority is not stolen,
- intervention does not escalate destructively.

Appeasement can create short-term stability,
but often strengthens misrecognition and dependency over time.

Scope × COS does not select appeasement as safety.

12.5 Minimum Preconditions for Use

Scope × COS does not require expertise.
But it does require minimal operational premises such as:

- an intention to keep judgment on the human side,
- acceptance that correctness will not be delivered,
- respect for the decision to suspend intervention.

If these premises do not hold, Scope × COS will not function as intended.

12.6 Domains Intentionally Out of Scope

Scope × COS intentionally does not address:

- therapy, healing, or recovery,
- moral evaluation or ethical instruction,
- measurement of personality or capability,
- optimization of relationships.

By excluding these domains, Scope × COS remains a structure for
making the limits of intervention visible.

12.7 Conclusion of This Chapter

Scope × COS does not claim universality or guaranteed effectiveness.

What it provides is:

> a reference point for identifying
> where intervention remains possible,
> and where it becomes destructive.

Scope × COS does not guarantee understanding or agreement.

However, it has a specific role:
to preserve the possibility of connection
even when understanding does not occur.

————————

Appendix A

COS v2 — Cognitive OS Internal Reference Model
(Operational Reference Excerpt)

This appendix presents the internal reference structure of
COS (Cognitive OS) as used within Scope × COS,
limited strictly to the minimum definitions required for operational reference.

COS is not a model for evaluation, classification, or diagnosis.
Nor is it a framework intended to fit individuals into predefined types.

Within the main body of Scope × COS,
COS is used solely as an "internal reference layer"
and never functions as the主体 of operational judgment or value determination.

■ Positioning of COS

COS is not a normative model that describes
how human cognition should be.

It is an internal reference model designed to structurally observe
which processing stages are active and what is occurring within them.

Scope defines the external frame for intervention, non-intervention, and stopping.
COS corresponds to the structural map of cognitive processing
that is referenced within that frame.

COS is not intended to be operated independently.
It is referenced only under the constraints defined by Scope.

■ Basic Structure of COS v2

COS v2 consists of the following primary layers:

A-Series (Premise Processing Layer)

C-Series (Output Control Layer)

D-Series (Observation, Control, and Synchronization Layer)

The B-Series exists as an internal processing layer.
However, within the operational context of Scope × COS,
it does not function as a direct basis for judgment or explanation,
and therefore is not foregrounded in this appendix.

■ A-Series (A1–A4): Premise and Decision Processing Layer

The A-Series is an internal layer
that handles how input is processed under specific premises.

A1: Premise Recognition

Determines under which premises incoming information is interpreted.
This is the most critical reference point within Scope × COS.

A2: Pathway Selection

Selects which cognitive pathway to proceed along
once the premises have been set.

A3: Possibility Expansion

Handles the expansion of possible trajectories
along the selected pathway.

A4: Implementation Judgment

Determines whether and how thought is translated
into action or decision.

* In Scope × COS,
  the outputs of the A-Series are not presented directly
  as external output.

■ C-Series: Output Control Layer

The C-Series does not control correctness of content.
It controls the form, intensity, and timing of output.

For example:

- whether to state something conclusively
- whether to return options
- whether to reduce output intensity
- or whether to remain silent

Such determinations are governed by the C-Series.

Within Scope × COS,
the C-Series is not used for persuasion, manipulation, or inducement.

■ D-Series: Observation, Control, and Synchronization Layer

The D-Series observes the overall cognitive process
and functions as a control layer
to prevent runaway processing or fixation of misrecognition.

Specifically, it detects signs such as:

- whether misreading or premise mismatch is becoming fixed
- whether the judgment authority is shifting externally
- whether intervention is beginning to become destructive

Within Scope × COS,
decisions regarding "stopping" or "termination of intervention"
are primarily based on observation by the D-Series.

■ Treatment of COS within Scope × COS

COS is not a completed map for understanding people.

Within Scope × COS, COS is used under the premises that it:

- does not evaluate
- does not classify
- does not produce conclusions

It is used solely as a reference structure
for identifying the limits of permissible intervention.

■ Scope of Public Disclosure of COS v2

COS v2 includes more detailed internal specifications.
However, at present,
the content provided in this appendix constitutes
the only reference information disclosed externally.

This appendix intentionally extracts only the portions
necessary to understand the positioning and operational role
of COS within Scope × COS.

While a full specification of COS v2 may be released separately in the future,
this document does not presuppose or depend on that release.

[Figure Note]

This figure is a conceptual support diagram intended to prevent misuse of the term “Layer” within the Scope × COS framework.
The layers shown here (L1–L8) do not represent value, superiority, or achievement goals.

- Purpose: Movement between layers involves cost, and such movement is not always symmetrical
  (difficulty of return and irreversibility may occur).
- The route depicted on the right does not indicate a recommended shortcut;
  it is a conceptual representation of a one-way transition that may occur under specific conditions.
- This figure is not a procedure manual or an operational guide.
  It must not be used on its own as a basis for classifying, evaluating, or directing people.

For operational definitions and constraints, refer to the main text and the appendix (COS).

Note:
This figure does not represent a cognitive hierarchy model itself,
nor does it suggest that reaching higher cognitive layers is inherently superior.
It serves solely as a supplementary visual aid
to intuitively convey transition load and the risk of misuse.

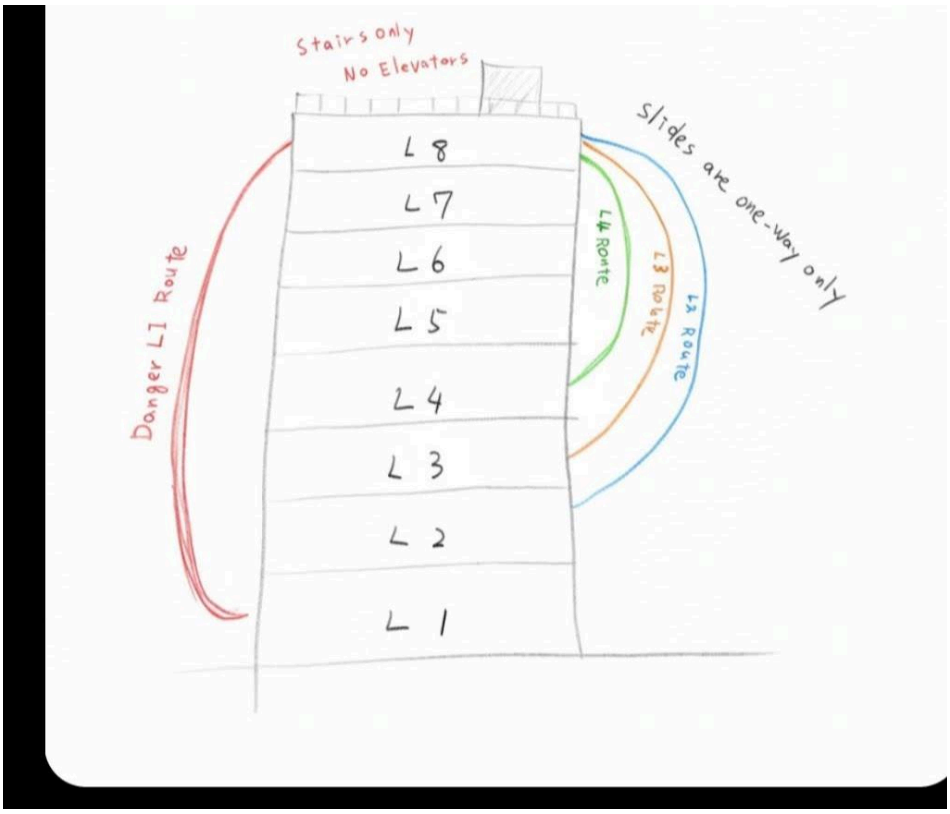

[Appendix Note]

This appendix constitutes internal reference information within Scope × COS.
All premises, constraints, and non-applicable scopes stated in the main text
also apply fully to this appendix.

The contents of this appendix are not intended to serve as a basis
for evaluation, classification, or judgment.
Operational judgment and responsibility are always retained on the human side.