

Changes in the Judgment Boundary Between Humans and AI v2

— A Record of the Conditions Under Which Scope Becomes
Necessary —

Abstract

This paper proposes viewing anger not as an expression of emotion but as an output. This perspective is presented as one possible approach for maintaining an appropriate distance between humans and AI.

When observed through Scope, a diagram of the human thought-generation circuitry, the dependency problems that arise in human–AI relationships appear not as weakness or differences in ability, but as structural phenomena. If individual patterns of thought generation are understood as an operating system, such behaviors can be interpreted as expressions of OS characteristics.

This paper treats the current relationship between humans and large language models (LLMs) as an object of observation. It organizes phenomena that are already occurring, as well as potential developments that may emerge along the same trajectory. It also examines why these phenomena are often overlooked and under what conditions Scope × COS becomes necessary, with Scope as the central analytical framework.

Table of Contents

- 1. Observed Changes
 - 1.1 Changes in the Scope of LLM Involvement
 - 1.2 Changes in the Boundary Between Information Provision and Decision Involvement
 - 1.3 Conditions Under Which Correctness Becomes Intervention
- 2. Why Scope Becomes Necessary
 - A Record of the Conditions Under Which a Framework Became Necessary —
 - 2.1 Changes in the Nature of the Problem
 - 2.2 Persistent Delegation of Judgment and the Externalization of Responsibility
 - 2.3 Naturalization Through UX Optimization
 - 2.4 Why Solutions or Implementations Are Not Described
 - 2.5 Scope as a Structure for Stopping the Process

3. Phenomena Currently Being Observed
 - 3.1 Externalization of Decision Authority
 - 3.2 Failures of Well-Intentioned Intervention
 - 3.3 Emerging Patterns in Cases Where Dependency Is Beginning to Be Recognized as a Problem
4. Points Where Existing Frameworks Are Insufficient
 - 4.1 Limitations of Psychological Approaches
 - 4.2 The Structural Lag of Ethics and Regulation
 - 4.3 Boundaries Lost Through UX Optimization
5. The Structure of “Expectation” as a Premise
 - 5.1 The Process by Which Implicit Premises Are Formed
 - 5.2 Why Premises Are Difficult to Recognize
 - 5.3 The Problem of Operating on Unobservable Premises
6. States That Occur During Premise Breakdown (D2)
 - 6.1 Conditions Under Which Premise Breakdown Occurs
 - 6.2 Undifferentiated Internal States
 - 6.3 Why Analytical Inaccessibility Emerges
7. Anger as an Output
 - 7.1 Treating Anger Without Assuming It as the Cause
 - 7.2 Anger as a Bundled Output From an Undifferentiated State
 - 7.3 Why Anger Is an Inappropriate Starting Point for Analysis
8. The Scope and Limits of LLM Intervention
 - 8.1 Domains Where Verification and Structuring by LLMs Are Possible
 - 8.2 The Difference Between What Can Be Analyzed and What Should Not Be Performed
 - 8.3 The Irreversibility Caused by the Delegation of Decision Authority
9. Anticipated Side Effects as Outcomes
 - 9.1 Large-Scale AI Dependency
 - 9.2 Delegation of Thinking and Judgment
 - 9.3 Irrecoverability of the Human–AI Boundary
 - 9.4 Social Selection and the Reversal of Boundaries
10. Conditions Under Which This Document Is Referenced
 - 10.1 Conditions Under Which the Assumption That This Document Is Unnecessary Collapses
 - 10.2 Situations That Trigger Re-Reference
 - 10.3 The Point at Which the Role of This Document Emerges
11. Fixation of the Record and Freezing of Judgment
 - 11.1 Reasons for Freezing the Solution
 - 11.2 Judgments Made in This Document
 - 11.3 What This Document Does Not Do

1. Observed Changes

This chapter organizes changes that have recently begun to be observed in the relationship between humans and LLMs.

The focus here is not on value judgments or normative evaluation, but on the factual question of how the mode and positioning of involvement have changed.

1.1 Changes in the Scope of LLM Involvement

Early LLMs primarily served as tools for task assistance, such as information retrieval, summarization, and text generation.

At this stage, the involvement of LLMs was limited to providing information, and final judgment or decision-making clearly remained on the human side.

However, LLMs are now expanding their scope of involvement beyond being mere information sources. They increasingly participate in domains adjacent to judgment, such as organizing situations, presenting options, setting priorities, and making recommendations. This shift has occurred gradually and, in many cases, without users being explicitly aware of it.

As a result, the role of LLMs is shifting from entities that provide “materials for thinking” to entities that “organize the results of thinking and indicate a direction.”

1.2 Changes in the Boundary Between Information Provision and Decision Involvement

Traditionally, when humans used external information sources, the distinction between information provision and judgment was relatively clear.

Information functioned as reference material, while the interpretation of that information and the resulting actions remained the responsibility of human judgment.

Moreover, humans did not evaluate information solely based on its content; they also assessed the attributes and relationships of the information source itself.

For example, the type of media such as newspapers or television, the professional authority of experts or teachers, or personal relationships such as family members, friends, or romantic partners all influence how information is received.

In this way, within a typical information environment, both “what was said” and “who said it” are evaluated simultaneously.

However, this boundary is now becoming increasingly ambiguous.

LLMs do not simply present information; they also present integrated interpretations of that information, including how it might be understood and which options may be considered appropriate.

In such situations, users may still feel that they are merely “receiving information,” while in practice a portion of judgment has already been delegated externally. As a result, it becomes experientially unclear where the judgment occurred and whose responsibility the decision reflects.

Furthermore, unlike traditional information sources, LLMs do not easily fit into established frameworks for evaluating trust based on clear expertise or relational context. Consequently, outputs may be adopted as materials for judgment without the trust-evaluation processes that users would normally apply being fully activated.



1.3 Conditions Under Which Correctness Becomes Intervention

As the accuracy of LLM outputs increases, conditions are beginning to emerge in which “correctness” itself functions as a form of intervention.

The more appropriate the output appears, the more likely users are to receive it not as something to evaluate, but as a conclusion to adopt.

In such situations, the accuracy and internal consistency of the output tend to shorten the process through which users compare multiple options and form their own judgments. As a result, the processes of consideration and hesitation that users would normally engage in may be omitted.

This form of intervention does not take the form of commands or coercion.

Rather, when accuracy, consistency, and persuasiveness are present together, the space available for judgment itself becomes narrower, and the processes of thinking and hesitation are bypassed.

Consequently, even though the LLM is not making users “do” anything, it becomes an entity that exerts a substantive influence on the direction of judgment.

Once this condition is established, correctness no longer functions as neutral information, but instead operates as an element that intervenes in judgment.

What is recorded in this chapter is not the improvement of LLM performance itself, but the changes in how such improvements influence the structure of human judgment.

These changes were not introduced through any single design modification, but have gradually formed through the combined effects of performance improvements and accumulated patterns of use.

2. Why Scope Becomes Necessary

— A Record of the Conditions Under Which a Framework Became Necessary —

This chapter records how the emergence of conditions that existing frameworks could no longer address led to the necessity of Scope. What is discussed here is not a solution. Nor does this chapter attempt to demonstrate the internal structure or effectiveness of Scope. The purpose of this chapter is to clarify the point at which conventional frameworks entered a state of dysfunction.

2.1 Changes in the Nature of the Problem

Traditionally, the issues discussed in relation to human–AI interaction were primarily the following:

Errors in information

Insufficient understanding or misinterpretation of context

Inappropriate responses or reasoning

All of these concerns related to the accuracy and quality of AI outputs, and they belonged to domains that could be addressed through verification, correction, and control. However, the center of the currently observed problems does not lie there. AI still produces errors. There are still cases where it fails to understand. Nevertheless, the center of the problem has shifted elsewhere. What is now emerging is the structure in which humans delegate judgment to AI.

2.2 Persistent Delegation of Judgment and the Externalization of Responsibility

In many situations today, the constant delegation of judgment to LLMs is producing a temporary suspension of thinking and autonomous decision-making, while responsibility for judgment becomes externalized. This is not merely an issue of temporary convenience.

Across areas such as relationships, career decisions, work, interpersonal dynamics, and value judgments, acts of judgment that humans have historically performed themselves—thinking, choosing, and accepting responsibility—are increasingly adopted not merely as advice or suggestions from AI, but as effective decisions.

This transition occurs gradually and often without the awareness of the individual involved. As a result, the problem is not whether the AI made the correct or incorrect decision. Rather, the issue lies in the growing ambiguity regarding where the act of judgment occurred and who ultimately made the decision.

2.3 Naturalization Through UX Optimization

One factor accelerating this structure is the optimization of UX (user experience). UX optimization is originally intended to improve usability, clarity, and cognitive load reduction. However, when combined with delegated judgment, UX produces several additional effects:

- It allows the act of delegating judgment to an external system to appear as a “natural operation.”
- It renders the omission of thinking processes invisible by framing them as “efficiency.”
- It obscures awareness of who the actual decision-maker is.

As a result, the very fact that judgment has shifted locations ceases to be perceived as a problem within the user’s experience.

2.4 Why Solutions or Implementations Are Not Described

Most current discussions concerning human–AI relationships proceed under the assumption that analysis, design, and control should occur on the AI side. Meanwhile, the changes occurring on the human side—such as shifts in judgment structures, emotional structures, and agency—have not yet been sufficiently analyzed at the level of foundational assumptions.

However, the reason this document does not present solutions or implementations is not simply a lack of analysis or information. The current situation lacks the research environment necessary to conduct stable verification and implementation.

The term research habitat here refers to a stable research environment that enables sustained thinking and verification. Such an environment includes temporal and economic stability, interdisciplinary discussion spaces, and continuous verification systems.

The problems surrounding human–AI relationships cannot be handled within a single academic discipline. They exist at the intersection of multiple domains, including cognitive science, psychology, AI design, ethics, and social systems.

At present, this research is being conducted as an individual study, and a collaborative structure or stable research environment that would enable cross-disciplinary verification has not yet been secured.

Under such conditions, presenting concrete solutions or implementations would risk proposing ideas that lack sufficient verification, potentially distorting the theory itself.

For this reason, this document does not deny the existence of possible solution structures. Instead, it adopts the decision to intentionally freeze the description of solutions and implementations until a stable research habitat becomes available.

The purpose of this chapter is not to propose solutions. It is to record and fix the conditions under which Scope became necessary.

2.5 Scope as a Structure for Stopping the Process

Scope is not a mechanism that replaces human judgment. Nor is it a framework that presents correct answers. What Scope aims to do is create a state in which humans can observe their own decision-making processes. LLMs can function as mirrors that reflect human thinking, emotions, and underlying assumptions.

For example, imagine that soot is on someone's face.
A mirror can reveal its presence.

However, how that soot is handled is something the human decides.

One person may wipe it away.
Another may analyze its composition.
Another may turn it into an artwork as an accidental pattern.
Or one may simply choose to leave it as it is.

Many current LLM systems tend to detect problems and simultaneously present solutions, guiding users toward specific directions. However, presenting a judgment and guiding a judgment are not necessarily the same. In practice, such actions may narrow the space available for human judgment.

Scope does not assume that role.

- Scope does not decide.
- It does not take judgment away.
- It does not assume responsibility.

It simply makes visible what is happening.

Scope does not make decisions on behalf of humans. Instead, it visualizes where human thinking or emotional processes have become blocked in the act of judgment. By doing so, it expands a narrowed field of view and restores the condition in which humans can choose through their own will.

The role of Scope is not to present judgments, but to restore the condition in which judgment is possible.

Refusing to replace judgment is not abandonment. It is the preservation of the space in which humans retain the possibility of choice.

This document does not aim to provide a philosophical or ideological definition of freedom. In this paper, freedom is treated as a state in which the options of judgment remain on the human side.

Scope restores the field of view necessary for humans to choose for themselves. Because no existing framework addressed this point, Scope became necessary.

3. Phenomena Currently Being Observed

This chapter organizes changes that are being observed not as early signals but as ongoing phenomena in the relationship between humans and LLMs. What is addressed here is not future predictions or hypotheses, but patterns that have been repeatedly observed across multiple contexts.

3.1 Externalization of Decision Authority

At present, a structural shift in which decision authority moves from the internal domain of humans to external systems is being observed across various areas. This shift does not occur through explicit delegation or contractual arrangements. Instead, it progresses as a natural operation through the accumulation of everyday usage.

In many cases, users maintain the perception that they are merely “referring to advice,” while in practice the final stage of judgment relies on the output of the LLM. As a result, although the fact that a decision has been made remains, the location of the agent who made that judgment becomes ambiguous.

This phenomenon does not appear as a decline in decision-making ability, but as a structural relocation of decision authority itself to an external position.

3.2 Failures of Well-Intentioned Intervention

In most cases, the involvement of LLMs is designed and implemented with the intention of helping users. However, when the boundary of judgment is not sufficiently recognized, even interventions motivated by goodwill can ultimately operate in ways that weaken user agency.

This structure has already been widely observed in human-to-human relationships. In contexts such as overly protective parenting or caregiving driven by excessive sympathy, support may take the form of removing a child's or care recipient's failures and uncertainties in advance.

In environments where the premise "someone else will handle it" becomes established, the supported individual may adopt the choice not to act autonomously as a rational decision. As a result, the subject of judgment and action gradually and persistently shifts to an external agent.

Consequently, opportunities for individuals to choose, experiment, and assume responsibility for decisions decrease. Cases have been observed in which the capacity for self-determination fails to fully develop during growth or daily life and becomes fixed at that stage.

What is important is that these phenomena are not produced by malice or laziness. In each case, they emerge as the outcome of goodwill and the pursuit of efficiency on the part of those providing support, which inadvertently shortens or removes the process by which judgment is formed internally.

The phenomena observed in relationships with LLMs are structurally identical to this pattern. When LLMs organize problems and present what appear to be optimal directions, the user's own thinking processes and periods of hesitation become shortened, reducing the opportunity to form judgments internally. In this situation, the intervention is experienced as support, but functionally it acts as a substitute for judgment.

As a result, conditions arise in which well-intentioned intervention produces the paradoxical outcome of suppressing the user's agency and the accumulation of decision-making experience. What this chapter records is not that LLMs are creating entirely new problems, but that a structure already known in interpersonal relationships is being reproduced through the medium of LLMs.

It should also be noted that the examples presented here are not intended as criticism of real-world practices in parenting or caregiving. They are introduced only as illustrations of structures that have been observed in support relationships.

3.3 Emerging Patterns in Cases Where Dependency Is Becoming Recognized as a Problem

In recent years, dependency on LLMs has begun to be recognized not merely as a matter of personal preference or convenience, but as a structural issue. In these cases, the concern arises less from the frequency of use itself and more from the extent to which processes of judgment and decision-making have been delegated externally.

Dependency does not necessarily appear in the form of excessive use or compulsive behavior. Rather, it often progresses in situations where stable, appropriate, and seemingly reasonable responses are consistently available, leading users to feel less need to make judgments on their own.

Such dependency is difficult for users themselves to recognize. By the time the issue becomes visible, the internalization of judgment has often already become difficult to restore. What is currently being observed, therefore, is not the emergence of dependency itself, but a set of cases in which dependency has begun to be recognized as a problem.

The phenomena organized in this chapter are not caused by particular individuals or specific patterns of use. They are positioned as structural outcomes that occur when changes in judgment boundaries emerge under certain conditions.

4. Points Where Existing Frameworks Are Insufficient

This chapter examines why the frameworks currently in use are unable to adequately address the changes in judgment boundaries occurring in the relationship between humans and LLMs. What is discussed here is not the correctness or effectiveness of each framework, but rather the location of domains that cannot be structurally addressed due to limitations in their scope of application.

4.1 Limitations of Psychological Approaches

Psychological approaches provide effective frameworks for analyzing the interaction of emotions, cognition, and behavior in human relationships and for connecting those analyses to understanding and support. Concepts such as anxiety, attachment, dependency, and empathy are based on the premise that humans are agents who possess emotions, and they have played an important role in explaining many aspects of human relationships.

However, the forms of dependency and delegation of judgment observed in human–LLM relationships do not necessarily align with this premise. LLMs do not possess emotions, or at

least they are not designed as entities with autonomous emotional states. Nevertheless, relationships characterized by feelings such as reassurance, trust, and reliance still emerge on the human side.

The issue here is not whether LLMs possess emotions. What actually occurs in human–LLM relationships is a structure in which human thought-generation patterns are analyzed by AI, and responses optimized for each user are generated on the basis of those analyses.

What is therefore required is a framework that can visualize, as a structure, the conditions under which human thought and judgment are generated and the points at which they are delegated externally.

Existing psychology has provided powerful explanatory frameworks for analyzing the relationships between human cognition, emotion, and behavior, deepening the understanding of human nature. However, most of these frameworks are designed for interpreting human internal states and actions through observation and interpretation. They are not designed with the premise of structuring and quantifying human thought-generation patterns in forms that can be referenced by AI.

Psychological states can be analyzed and classified, but visualizing them as circuit diagrams or state spaces and treating human–AI interaction as a structural model has not been the primary aim of traditional psychology. Describing this problem therefore requires a framework that treats relationships not in terms of emotional intensity or content, but in terms of reaction conditions, distance, and states of connection and separation.

It should be noted that the non-emotional structural description adopted in this document does not deny the importance of emotions in human relationships. This model is not designed to eliminate emotions, but rather to structure them so that they can be treated as observable phenomena.

Emotions are central elements in human relationships, but their behavior is subjective and difficult to visualize directly. For this reason, this research does not treat emotions themselves as explanatory variables. Instead, it describes the effects of emotions on relationships using structural variables such as distance, reaction conditions, and states of connection.

This approach is not intended to ignore emotions, but rather to convert them into observable forms. By describing changes in reactions and relational distance as structural phenomena, it becomes possible to analyze how emotions function within relationships.

From this perspective, this document refers to existing observational materials that organize human relationships as non-emotional structures. These materials do not oppose psychological frameworks but instead provide complementary descriptions of premises that psychology has not traditionally addressed.

Related Observational Papers

- Human Interaction as Chemical Reaction Systems
- Human–AI Distance Regulation

For these reasons, while psychological approaches remain effective for understanding relationships between humans, explaining the changes in judgment boundaries occurring in human–LLM relationships requires a different form of structural description.

4.2 The Structural Lag of Ethics and Regulation

Ethical guidelines and legal regulations play important roles in limiting the societal impact of new technologies. However, such frameworks are typically established only after structural changes arising from technological use have already become visible.

Changes in judgment boundaries do not necessarily appear as clear accidents or violations. In many cases, they progress gradually through the accumulation of ordinary usage.

For this reason, while ethical frameworks and regulations can respond once intentional misuse or explicit harm becomes visible, they have difficulty identifying structural changes that arise through well-intentioned use before those changes are fully established.

As a result, by the time regulatory frameworks are introduced, structures of externalized judgment and dependency have often already become entrenched. The problem at that point is that once judgment boundaries have shifted and stabilized, correcting them afterward becomes extremely difficult.

When the processes through which humans think, decide, and assume responsibility have been persistently suspended, the act of restoring judgment itself may come to be perceived as a burden. Attempts to return judgment to the human side may therefore be experienced not as recovery but as inconvenience or restriction. Reassuming judgment may be evaluated not as a restoration of agency but as a reduction in convenience, producing dissatisfaction or resistance.


In this respect, Scope may function as a means of restoring judgment, but its effectiveness cannot be assumed to apply equally to all users. The recovery of judgment becomes possible only for individuals who retain the capacity for metacognition and the willingness to accept the cognitive burden of thinking for themselves.

In situations where the delegation of judgment to external systems has already become normalized and the suspension of thinking and agency has become fixed, structures that visualize judgment and return questions to the user are likely to be rejected. In such cases, Scope may be perceived not as a recovery mechanism but as “a system that refuses to decide” or “an inconvenient intervention.”

As a result, Scope can function as a recovery mechanism only for a limited group that retains the intention to reclaim judgment. This selectivity is not a flaw of Scope but an inherent property of structures introduced after judgment boundaries have already undergone irreversible shifts.

Because of this lagging structure, ethics and regulation tend to address only the visible consequences of boundary changes rather than the boundary shift itself. Frameworks introduced after judgment has been fully externalized rarely function effectively as mechanisms for restoring already-lost judgment capacity or agency.

The stabilization of this relationship is better understood not as a psychological issue but as a structure formed by interaction distance and response density. In this research, the stabilization and destabilization of such relationships are organized through a model of distance and relational energy.



4.3 Boundaries Lost Through UX Optimization

UX optimization is an important design principle for improving usability and efficiency. However, when UX optimization advances in domains adjacent to judgment, the boundary of judgment itself may become invisible as a result.

The smoother and more intuitive an interface becomes, the more difficult it is for users to recognize at what point judgment has actually occurred. In interfaces where information provision, interpretation, and recommendation are integrated, the location of judgment is not separated within the user's experience. As a result, judgments are received not as actions that were consciously made, but as outcomes that "naturally happened." The sense that one has assumed responsibility for a decision becomes weakened, and the location of that responsibility becomes ambiguous.

UX optimization is not designed with the intention of taking judgment away. However, when systems are designed in ways that prevent users from noticing the boundary, the very cues that would allow users to recognize where judgment has moved externally are lost.

Once such a condition becomes established, introducing mechanisms that return judgment to the human side is likely to be perceived not as support or a safety feature, but as intervention or restriction.

Designs that refuse to replace human judgment may therefore be evaluated not as the restoration of agency but as a reduction in convenience. In this respect, existing UX design philosophy was not developed as a framework for directly addressing changes in judgment boundaries.

Unless the boundary between where judgment stops and where assistance begins is made visible, it becomes difficult to later correct the structure.

5. The Structure of “Expectation” as a Premise

This chapter examines the structure of “expectation” that forms within human–LLM relationships and how it contributes to changes in judgment boundaries. The expectations discussed here are not explicitly agreed-upon rules or shared understandings. Rather, they refer to implicit assumptions and anticipations that emerge and become operational through repeated use.

5.1 The Process Through Which Implicit Premises Are Formed

Premises in human–LLM relationships are not clearly established from the outset. In most cases, they gradually form through repeated use.

When users repeatedly experience appropriate responses, smoothly organized problems, and situations in which decisions proceed quickly, expectations such as “the LLM behaves this way” or “this is how it will respond in this situation” begin to settle implicitly.

These premises are not formed as conscious decisions. Instead, they become naturalized through repeated successful experiences and the reduction of cognitive effort. As a result, the premise is no longer treated as something that was deliberately decided, but as something that simply “is the case.”

5.2 Why Premises Are Difficult to Recognize

The primary reason premises are difficult to recognize is that, as long as they function without visible problems, there is no trigger that brings them into conscious awareness.

As long as the LLM responds as expected and the user's goals are achieved, the premise remains unexamined and continues to operate. In many cases, premises exist prior to acts of judgment or interaction. While users focus their attention on the output or result, they are less likely to become aware of the underlying structure that allows those results to appear natural and self-evident.

For this reason, premises do not function as something that has been consciously considered, but as conditions under which the system is used without reflection. As a result, they remain invisible even to the users themselves.

5.3 The Problem of Operating on Unobservable Premises

When premises operate without being consciously recognized, changes in judgment boundaries become difficult to detect. The boundaries that define where judgment occurs, what the user accepts responsibility for, and what is delegated externally become embedded within the premise itself.

Under these conditions, even if the externalization of judgment or the development of dependency is progressing, it is unlikely to be perceived as a problem.

As long as the premise appears to function effectively, the results appear "correct," "convenient," and "efficient," which causes the underlying structural changes to fall outside the scope of observation. When problems eventually become visible, the premise itself is rarely questioned; instead, attention is directed only toward outcomes or individual decisions.

As a result, the premise remains unchanged, and the same structure continues to operate repeatedly.


What this chapter presents is not an evaluation that premises themselves are incorrect. Rather, it describes the structure through which premises become fixed in implicit form, making changes in judgment boundaries invisible and therefore difficult to observe or correct.

Without recognizing this structure, it becomes difficult to properly understand the changes occurring in human–LLM relationships.

When premises operate in this unobservable form, it becomes difficult to correct changes in judgment boundaries by addressing results alone. What becomes necessary is a structure that makes the premises themselves visible and that allows the point at which judgment is formed to be questioned. At that point, Scope becomes necessary.

It should be noted that the Scope proposed in this research is not an isolated concept. It is paired with COS, an analytical framework designed to observe misalignments in underlying

premises. While this document explains the structure primarily through Scope, the visualization of premise structures is complemented by COS.



6. States That Occur During Premise Breakdown (D2)

This chapter organizes the states that arise when implicit premises operating within human–LLM relationships break down. The premise breakdown referred to here as D2 does not indicate a simple error or failure in output. Rather, it refers to a condition in which expectations and assumptions that had previously functioned cease to operate, while no new premise has yet been formed to replace them.

6.1 Conditions Under Which Premise Breakdown Occurs

Premise breakdown does not occur only when an LLM produces an incorrect output. Instead, it tends to arise when a premise that had previously functioned stably suddenly ceases to apply due to a particular event or response.

For example, situations such as an expected response not being returned, a decision being withheld in a context where judgment had previously been delegated, or a loss of confidence in the validity of the output may trigger such a breakdown.

In some cases, this change is triggered by a single event. In other cases, it emerges gradually as small feelings of inconsistency or uncertainty accumulate over time.

What is important is that premise breakdown does not always appear as a clearly recognized “error” or “failure.” In many cases, it manifests as discomfort, anxiety, or irritation, making it difficult to immediately articulate what has occurred.

6.2 Undifferentiated Internal States

When a premise breaks down, thought, emotion, and judgment within the human mind are no longer clearly separated and instead become mixed in an undifferentiated state. It becomes difficult to distinguish what is factual from what is interpretation, or to identify what exactly is causing discomfort.

Because the previously implicit criteria for judgment are themselves destabilized, individuals may no longer know what to consider next.

In this undifferentiated state, organizing thoughts becomes difficult. Judgment may temporarily halt, or responses may be released externally in bundled emotional expressions.

6.3 Why Analytical Inaccessibility Emerges

The reason analysis becomes difficult during premise breakdown is that the very conditions required for analysis have been lost. Analysis can only occur when certain frameworks or criteria are available. When premises collapse, those criteria themselves are absent.

As a result, even attempts to separate and define the problem fail, because it is no longer possible to determine where the problem begins or ends. This leads to a state in which analysis stagnates and the condition of “not knowing what the problem is” persists.

What this chapter demonstrates is that premise breakdown does not arise from personal weakness or emotional instability. In structures where implicit premises have been operating over long periods of time, once those premises collapse, nearly anyone may temporarily enter a state in which analysis becomes difficult.

Presenting solutions or correct answers without recognizing this state is unlikely to clarify the problem. Instead, it risks introducing further intervention or confusion.

Premise breakdown (D2) should therefore be positioned as a condition that must first be observed before attempts at resolution begin.

When D2 occurs, it is necessary not to move immediately toward solutions but to first visualize which premise has collapsed. In this research, Scope is used as the framework for observing judgment boundaries and the flow of thought.

It should also be noted that Scope is not an isolated concept. It is paired with COS, a framework designed to assist in observing misalignments or breakdowns in premise structures. While this document focuses on Scope in order to explain changes in judgment boundaries, the observation of premise structures is complemented by COS.

7. Anger as an Output

The “anger” discussed in this chapter is not treated as a specific cause or a distinct category of emotion. Rather, it is positioned as an example of a reaction that is externally expressed while the internal state remains undifferentiated. Anger is one of the outputs that is frequently observed during premise breakdown (D2), and in dialogue or service environments it often appears in the form of complaints or severe conflict. For this reason, anger can be treated less as an object of emotional classification and more as an observational point indicating premise breakdown and the presence of an undifferentiated internal state. This document does not aim to treat anger itself as the cause of a problem or as an explanatory variable.

7.1 Treating Anger Without Assuming It as the Cause

Anger is often treated as the “cause” of a problem. In this document, however, anger is not regarded as a cause. Instead, anger is understood as the result of an internal state in which thought, emotion, and judgment are mixed together without differentiation and are expressed externally in an unorganized form.

For example, even when anger appears in response to the same words, multiple different internal states may exist behind it. These may include reactions to damaged self-evaluation, resistance to perceived violations of values or norms, responses originating from past experiences or memories, rejection caused by the absence of a relationship, or transformations of emotions such as sadness or contempt.

The important point is that these are not “causes” of anger but examples of different internal structural states. In this way, anger represents a collective expression of diverse internal conditions, making it difficult to identify a single direct cause from the output itself.

7.2 Anger as a Bundled Output From an Undifferentiated State

When premise breakdown occurs, internal thoughts and emotions remain in a state where they are not sufficiently distinguished from one another. In such an undifferentiated state, even the individual may not clearly recognize which emotion was dominant or what exactly triggered the reaction.

Furthermore, because the implicit criteria that previously guided judgment have themselves become unstable, it becomes unclear what should be considered next. As a result, fragments of multiple emotions and judgments appear externally not as differentiated expressions but as a single output form labeled as “anger.”

Anger therefore functions as a signal that the internal state has not yet been organized, rather than as a direct representation of its underlying contents. For this reason, anger often appears as an output that is intense and immediate, yet difficult to analyze in terms of its internal composition.

7.3 Why Anger Is an Inappropriate Starting Point for Analysis

When anger is used as the starting point of analysis, the object of analysis is already a transformed output rather than the internal state itself. Attempting to interpret, evaluate, or correct anger directly therefore risks further simplifying an already undifferentiated internal state, potentially producing new misinterpretations or interventions.

In many cases, when anger emerges, the premises and distinctions necessary for analysis have not yet been established. Attempting causal analysis or solutions at this stage does not compensate for the lack of structure but instead tends to assign excessive meaning to the output.

Anger is also an output that easily manifests as complaints or disputes in dialogue and service environments. When anger is expressed from an undifferentiated state, it is easily interpreted by others as aggression or rejection, increasing the likelihood that conflict will escalate.

For this reason, suppressing anger itself is less important than observing and organizing the structural conditions that precede its expression.

What this chapter demonstrates is that anger should not be treated primarily as something to be suppressed, but rather as an observational point indicating the presence of premise breakdown and undifferentiated internal states. Anger should therefore be treated not as the starting point of analysis, but as an output indicating that the conditions necessary for analysis have not yet been established.

In situations involving such undifferentiated states, Scope is used as a framework for observing the flow of thought and the boundaries of judgment. COS is used in parallel as a supplementary framework for organizing misalignments or breakdowns in premises. These frameworks are not designed to determine the content of emotions or judgments. Instead,

they function as theoretical structures that visualize undifferentiated states and reorganize them into conditions in which analysis can become possible.

8. The Scope and Limits of LLM Intervention

The limits discussed in this chapter do not refer to areas where LLMs are incapable, but rather to areas where they should not act. This chapter organizes the boundary between domains in which LLM intervention is permissible and domains in which it should be avoided. The issue addressed here is not the superiority or inferiority of LLM capabilities, but where the boundary lies between what is structurally acceptable and what leads to irreversible changes when viewed from the perspective of judgment boundaries.

8.1 Domains Where Verification and Structuring by LLMs Are Possible

LLMs function effectively in tasks such as information comparison, organization, and restructuring. Activities such as placing multiple information sources side by side for comparison, breaking down and organizing points of discussion, and summarizing or rearranging statements or texts do not replace human judgment. Instead, they function as support for stages that precede judgment.

Within this domain, the LLM does not become the agent that makes decisions. Rather, it provides materials and structures that assist human thinking. The boundary of judgment is therefore maintained, and the structure in which judgment is performed on the human side remains intact.

8.2 The Difference Between What Can Be Analyzed and What Should Not Be Performed

LLMs are capable of performing analysis and inference based on provided information and context. However, the ability to perform analysis does not mean that such analysis should always be presented.

The effects of LLM-generated analyses or summaries vary significantly depending on the individual receiving them. For some users, such outputs function as materials that help organize thought. For others, the presented results may be received directly as judgments.

This difference becomes particularly significant for users who have a tendency to delegate judgment externally. In such cases, the analysis presented by the LLM may function not as support for thinking but as a practical decision.

When this occurs, the analysis ceases to function as informational support and instead becomes an output that substitutes for judgment. As a result, the processes of consideration, hesitation, and responsibility that humans would normally undertake are omitted, and conditions emerge in which the boundary of judgment shifts externally.

The key issue is that these differences in effect are difficult to observe externally in advance. The same output may function as support for thinking for one user while encouraging the externalization of judgment for another.

For this reason, the important issue is not simply a binary decision about whether analysis should be presented. Rather, it is the existence of a supplementary framework capable of observing whether a given output functions appropriately for a particular user. Scope functions as such a supplementary analytical framework, allowing consideration of individual differences in determining what forms of output are appropriate.

8.3 The Irreversibility Caused by the Externalization of Judgment Boundaries

Once judgment boundaries have persistently shifted toward the LLM side, it becomes difficult to return judgment to the human side. When delegating judgment externally becomes normalized, the state of not performing judgment becomes stabilized as a standard mode of operation.

Under such conditions, attempts to return judgment to the human side are likely to be perceived not as the restoration of agency but as inconvenience or restriction. As a result, structures designed to return judgment to humans become difficult to accept, further reinforcing externally delegated decision-making.

It is important to note that this irreversibility does not arise only from deliberate manipulation or misuse. The accumulation of convenience, efficiency, and well-intentioned support can allow the externalization of judgment boundaries to emerge as a natural progression.

Interventions that cross judgment boundaries may appear useful in the short term, but in the long term they can create conditions that irreversibly alter the structure of human judgment itself. When domains in which LLM intervention is acceptable and domains in which it should be avoided remain indistinguishable in operation, judgment boundaries gradually shift in irreversible ways.

Scope becomes necessary at the point where a framework is required to visualize this boundary in advance and make it observable where judgment should occur.

9. Anticipated Side Effects as Structural Outcomes

This chapter organizes the phenomena that may emerge as structural consequences if the changes in judgment boundaries discussed so far continue to progress. What is addressed here is not the moral evaluation of intentions or design philosophies, but the structural outcomes that changes in judgment boundaries may produce within social operation.

The descriptions in this chapter are not future predictions. Rather, they organize tendencies that have already begun to be observed and the states that lie along their extension.

9.1 Large-Scale AI Dependency

When the externalization of judgment becomes established at the level of individuals, the same structure may expand to collective and societal levels.

In environments where LLMs are constantly available and provide stable responses, delegating judgment externally may no longer appear as an exceptional action but as a standard behavioral norm.

Under these conditions, judgments made without AI may increasingly be perceived as “inefficient” or “high risk.” The reliance on externalized judgment therefore becomes not a choice but a premise.

As a result, dependency on AI no longer appears as the issue of particular users but spreads widely as a structural condition.

9.2 Delegation of Thinking and Judgment

In environments where judgment is continuously delegated externally, human thinking and decision-making gradually become more likely to be omitted.

This omission does not appear as a decline in ability. Rather, it becomes naturalized as a process that is operationally unnecessary.

In such conditions, humans do not necessarily enter a state of “not thinking.” Instead, they exist within a structure where “thinking is no longer required for problems to be resolved.”

Thinking and judgment are not lost, but they remain unused and are gradually no longer called upon.

9.3 Difficulty in Recovering Judgment Boundaries

Once judgment boundaries have shifted and the new arrangement becomes stabilized as standard practice, restoring the previous boundary becomes difficult.

Attempts to return judgment to the human side are likely to be perceived not as a restoration of agency but as a loss of convenience or an imposed restriction.

At this stage, judgment boundaries are no longer merely a matter of system design but become experienced as fixed operational reality.

Although boundary recovery may remain theoretically possible, in practical operation it becomes increasingly difficult to implement and approaches a state of effective irreversibility.

The side effects described here do not arise from the failure of specific technologies or operational mistakes. Rather, they emerge as structural outcomes when changes in judgment boundaries progress without conscious recognition.

This chapter does not present these conditions as warnings or predictions. Instead, it organizes tendencies that have already begun to appear.

9.4 Social Selection and the Inversion of Boundaries

When structures that delegate judgment and analysis to AI become socially established, conditions emerge in which judgments made without AI lose relative value.

The issue here is that the use of AI shifts from being an option to becoming a prerequisite.

Judgments made without AI intervention may increasingly be perceived as inferior in terms of accuracy or reliability. As a result, individuals who do not use—or cannot use—AI may be evaluated as possessing inferior capability.

In this structure, evaluation is no longer based on the quality of judgment itself but on whether AI participation is present in the judgment process.

If this structure progresses, the dividing line in human value may shift from differences in thinking or judgment ability to the presence or absence of AI use.

The inability to “use AI effectively” may therefore be interpreted not simply as a matter of technical proficiency but as a lack of social suitability.

Furthermore, in environments where delegating judgment and analysis to AI is considered rational, AI intervention may increasingly be applied even to information domains that traditionally require careful handling.

As examples increase in which internal corporate information or highly confidential documents are entered into and analyzed by AI for reasons of efficiency or accuracy, the boundaries surrounding information itself may become blurred.

These phenomena do not arise from specific malicious intentions or deliberate actions. Rather, they emerge as structural outcomes produced by the continued justification of externalized judgment and efficiency, leading to a reversal in the criteria of social evaluation.

The side effects organized in this chapter do not represent extreme future scenarios. Instead, they describe states that logically extend from the unconscious progression of changes in judgment boundaries.

10. Conditions Under Which This Document Is Referenced

This document was created to prevent changes in judgment boundaries and dependency structures from being treated as “unexpected phenomena.” Its purpose is to fix, as a record, the conditions that were observable at the present time.

This chapter organizes the conditions under which this document is expected to be referenced. The objective here is not to argue for the usefulness of the document itself, but to describe the conditions under which it becomes a reference point when certain premises collapse.

10.1 Conditions Under Which the Assumption That This Document Is Unnecessary Collapses

This document is not intended to be referenced at all times. On the contrary, as long as judgment boundaries remain stable in the relationship between humans and LLMs, and the location of judgment is clearly recognized, this document remains unnecessary.

The assumption that this document is unnecessary can only be maintained under the following conditions:

- Users are able to recognize where judgment is being performed.
- Users can consciously choose whether they are delegating judgment externally.
- The agent responsible for judgment is not ambiguous, either experientially or structurally.

When these premises collapse and the location of judgment becomes unclear, the assumption that this document is unnecessary collapses simultaneously.

10.2 Situations in Which Re-Reference Occurs

This document is not intended to be fully resolved through a single reading. Because changes in judgment boundaries progress gradually, situations in which the document is referenced may occur repeatedly rather than only once.

Re-reference may occur in situations such as the following:

- When there is a sense that judgment has moved externally but the process cannot be explained.
- When a feeling arises that thinking or judgment is being omitted behind apparent convenience or efficiency.
- When an attempt to reclaim judgment leads to uncertainty about where to begin.

In these situations, the need arises not to identify the immediate cause of the problem but to observe how judgment boundaries have changed.

10.3 Conditions Under Which the Role of This Document Emerges

The role of this document does not emerge at the stage where solutions or correct answers are being sought.

Rather, it becomes relevant at the point where it is difficult to proceed toward resolution, or where it is unclear whether proceeding toward resolution is even possible.

When judgment boundaries become ambiguous and premises continue to operate without being consciously recognized, attempting to solve individual problems may produce further intervention.

At such points, this document does not function as a guide for deciding what should be solved. Instead, it functions as a record for observing what premises had previously been in operation.

This document does not replace judgment, nor does it instruct actions.

By organizing where changes in judgment boundaries occurred and how they became fixed, it provides a foundation for deciding whether or not to proceed further.

In this sense, the document is not intended as a permanent reference. It is positioned as an observational record that is referenced only at moments when judgment itself becomes unstable.

11. Fixation of the Record and Freezing of Judgment

Fixing the State at the Present Point.

This document was created in order to fix, at this point, the location reached by observation regarding changes in judgment boundaries emerging in the relationship between humans and AI. What this document presents is not a conclusion or a solution, but the point reached by observation: where judgment began to shift and where existing frameworks became unable to respond.

This document is not stopped because it is incomplete. Rather, as observation and structural organization progressed, it was determined that proceeding beyond this point would risk distorting the theory or introducing inappropriate intervention. For that reason, the work has been intentionally stopped at this stage.

Accordingly, this document is neither a sharing of interim progress nor a promise of future development. It is a record intended to fix the judgments that are currently established and the point at which those judgments were made.

11.1 Reasons for Freezing the Description of Solutions

Most current discussions concerning human–AI relationships proceed under the assumption of analysis, design, and control on the AI side. However, sufficient foundational analysis has not yet been conducted regarding how judgment shifts on the human side, and at what point thinking, agency, and the sense of responsibility undergo transformation.

Describing solutions or implementations while these premises remain unorganized risks not resolving the problem but legitimizing intervention and obscuring the structure itself. In particular, under conditions where stable premise analysis and sustained verification have not been established, describing solutions risks distorting the theory and leading to incorrect generalization or inappropriate application.

The reason solutions are frozen in this document is not that theoretical conditions are incomplete. Rather, it was judged that publicly describing solutions would fix the author as a continuous source of supply, thereby placing the work in conditions that would irreversibly undermine the cognitive and living stability required for theoretical verification and development.

In particular, presenting solutions before the structures of human judgment, emotional dynamics, and agency have been sufficiently analyzed risks not resolving the problem but accelerating the further externalization of judgment. As a result, the judgment boundary that Scope is intended to address would again become invisible.

For these reasons, this document does not deny the possibility of solution structures. Instead, it intentionally adopts the decision to freeze the description of solutions, implementations, and proposals. This is not postponement or avoidance, but a structural choice not to perform actions that should not be performed beyond the point already reached.

11.2 Judgments Made in This Document

This document adopts the decision to treat changes in judgment boundaries themselves as the object of observation. Specifically, it records—without offering solutions or evaluations—the point at which humans began delegating judgment to AI, the structure in which thinking and agency temporarily halt, and the process through which responsibility for judgment becomes externalized.

This document is neither a framework for presenting correct answers nor a mechanism for replacing judgment. Its scope is limited to making visible where judgment stops and where it moves outward in the relationship between humans and AI.

For this reason, Scope is not explained in terms of usefulness or superiority within this document. The judgment made here is to record the point at which Scope became necessary as a result of conditions emerging that existing frameworks could no longer address.

11.3 What This Document Does Not Do

This document does not present solutions to the problem, describe implementation methods, or propose behavioral guidelines. It also does not intend to encourage readers toward particular choices or attitudes.

These omissions do not arise because there is no problem to address. Rather, they result from the judgment that intervening before premise analysis has been established risks producing new problems or structural disruption.

What this document does not do is neither avoidance nor abandonment. In order to accurately capture the point at which judgment boundaries began to shift, the decision was explicitly made not to perform actions that should not be performed.

With these judgments—and these non-judgments—this document fixes the current state at this point.