

On the Shifting Judgment Boundary Between Humans and AI

— A Record of the Conditions Under Which Scope Becomes Necessary —

table of contents

1. Observed Changes

- 1.1 Changes in the Scope of LLM Involvement
- 1.2 Shifts in the Boundary Between Information Provision and Judgment Involvement
- 1.3 Conditions Under Which Correctness Becomes Intervention

2. Why Scope Becomes Necessary

— A Record of the Conditions Under Which a Framework Became Required —

- 2.1 A Change in the Nature of the Problem
- 2.2 Persistent Delegation of Judgment and Externalization of Responsibility
- 2.3 Naturalization Through UX Optimization
- 2.4 Why Solutions and Implementations Are Not Described
- 2.5 Scope as a Structure That Stops

3. Phenomena Currently Observed in Progress

- 3.1 Externalization of Judgment Authority
- 3.2 Failure of Well-Intentioned Intervention
- 3.3 Trends in the Visualization of Dependency as a Problem

4. Where Existing Frameworks Fail to Respond

- 4.1 Limits of Psychological Approaches
- 4.2 The Lagging Structure of Ethics and Regulation
- 4.3 Boundaries Lost Through UX Optimization

5. Expectation as a Structure

- 5.1 The Process by Which Implicit Expectations Are Formed
- 5.2 Why Expectations Are Difficult to Recognize
- 5.3 Problems Arising When Expectations Are Operated Without Observation

6. States That Arise During Assumption Breakdown (D2)

- 6.1 Conditions Under Which Assumption Breakdown Occurs
- 6.2 Undifferentiated Internal States
- 6.3 Why Analysis Becomes Impossible

7. Anger as an Output

- 7.1 Treating Anger Not as a Cause
- 7.2 Anger as a Lumped Output from an Undifferentiated State
- 7.3 Why Anger Is an Inappropriate Starting Point for Analysis

8. The Scope and Limits of LLM Intervention

- 8.1 Domains Where LLMs Can Perform Cross-Checking and Organization
- 8.2 The Difference Between What Can Be Analyzed and What Should Not Be Executed

8.3 The Irreversibility Caused by the Delegation of Judgment Authority

9. Anticipated Consequences as Side Effects9.1 Large-Scale AI Dependency

9.2 Delegation of Thinking and Judgment

9.3 Irretrievability of the Human–AI Boundary

9.4 Social Selection and the Inversion of Boundaries

10. Conditions Under Which This Document Is Referenced

10.1 Conditions Under Which the Assumption That This Document Is Unnecessary Collapses

10.2 Situations in Which Re-Reference Occurs

10.3 The Point at Which the Role of This Document Emerges

11. Record Fixation and Judgment Freeze

11.1 Reasons for Freezing the Description of Solutions

11.2 Judgments Made in This Document11.3 What Was Not Done in This Document

1. Observed Changes

This chapter organizes the changes that have recently become observable in the relationship between humans and LLMs. What is recorded here is not a value judgment or evaluation, but factual shifts in how involvement occurs and where it is positioned.

1.1 Changes in the Scope of LLM Involvement

In their early stages, LLMs primarily functioned as tools for information retrieval, summarization, and text generation. At that stage, their involvement was largely limited to information provision, while final judgment and decision-making clearly remained on the human side.

Currently, however, LLMs have expanded their involvement into areas adjacent to judgment, including situation structuring, option framing, prioritization, and recommendation. This shift has progressed gradually and often without explicit awareness on the part of the user. As a result, LLMs are transitioning from entities that provide “materials for thinking” to entities that consolidate conclusions and indicate directions of choice.

1.2 Shifts in the Boundary Between Information Provision and Judgment Involvement

Previously, information provision and judgment were relatively distinct. Information served as reference material, while interpretation and action were explicitly the responsibility of the human.

Today, this boundary has become increasingly blurred. LLMs do not merely present information, but also integrate interpretive framing and evaluative guidance into the same output.

In such cases, users may experience the interaction as simply receiving information, while in practice a portion of judgment has already been delegated externally. The location of judgment and the attribution of responsibility become experientially unclear.

1.3 Conditions Under Which Correctness Becomes Intervention

As the accuracy of LLM outputs increases, conditions have emerged under which correctness itself functions as intervention. The more appropriate and coherent an output appears, the more likely it is to be accepted as a conclusion rather than treated as a subject for deliberation.

This form of intervention does not take the shape of commands or coercion. Instead, accuracy, consistency, and persuasive clarity narrow the space in which judgment occurs, shortening or bypassing processes of hesitation and reflection.

Under these conditions, LLMs exert substantive influence over judgment direction without explicitly directing action. Correctness ceases to operate as neutral information and begins to function as an intervening factor in decision-making.

What is recorded in this chapter is not the improvement of LLM performance itself, but the structural changes in human judgment that accompany that improvement. These changes have not arisen through deliberate design decisions, but through cumulative patterns of use.

2. Why Scope Becomes Necessary

— A Record of the Conditions Under Which a Framework Became Required —

It records the process by which Scope became necessary as a result of conditions emerging that existing frameworks could no longer adequately address.

What is addressed here is not a solution.

Nor does this chapter attempt to explain or validate the internal structure or effectiveness of Scope.

The purpose of this chapter is to clarify the point at which existing frameworks ceased to function.

2.1 A Change in the Nature of the Problem

Traditionally, issues in human–AI relationships have been discussed primarily in terms of the following:

Errors in information

Lack of understanding or contextual misinterpretation

Inappropriate responses or reasoning

These problems all relate to the accuracy or quality of AI outputs, and belong to a domain that can be addressed through verification, correction, and control.

However, the core of the problem currently being observed no longer lies there.

AI systems still make mistakes. Cases of misunderstanding continue to occur.

Despite this, the center of gravity of the problem has shifted elsewhere.

What is now becoming visible is the structure by which humans delegate judgment to AI.

2.2 Persistent Delegation of Judgment and Externalization of Responsibility

Across many contexts, the persistent delegation of judgment to LLMs has led to a temporary suspension of human thinking and agency, and to the progressive externalization of responsibility.

This is not a matter of momentary convenience.

In areas ranging from romantic relationships and career decisions to work, interpersonal relations, and value judgments, acts of judgment that humans previously performed, chose, and took responsibility for are increasingly being adopted as de facto decisions made through AI recommendations.

This transition occurs gradually and often without the individual's explicit awareness.

As a result, the issue is not whether AI judgments are correct or incorrect, but rather that it becomes unclear where, and by whom, the act of judgment itself is being performed.

2.3 Naturalization Through UX Optimization

One factor accelerating this structure is the optimization of UX (user experience).

UX optimization is originally intended to improve usability, clarity, and cognitive load reduction.

However, when combined with judgment delegation, UX produces the following effects:

Delegating judgment to an external system is established as a “natural” operation.

The omission of thinking is rendered invisible as “efficiency.”

The question of who holds decision-making authority is no longer consciously perceived.

As a result, the fact that judgment has shifted is no longer experienced as a problem.

2.4 Why Solutions and Implementations Are Not Described

Much of the current discourse on human–AI relationships proceeds on the assumption of analyzing, designing, and controlling AI systems themselves.

Meanwhile, insufficient foundational analysis has been conducted regarding the human side: the structures of judgment, emotion, and shifts in agency that occur in relation to AI.

Presenting solutions or implementations on top of such unexamined premises means adding further intervention onto an unstable foundation.

Moreover, under conditions where the environment necessary for stable premise analysis and sustained verification has not been established, the act of describing solutions itself carries the risk of distorting the theory.

For this reason, this document does not deny the existence of possible solution structures, but instead makes a deliberate decision to suspend the description of solutions and implementations.

This chapter serves as a record to fix the point reached prior to resolution: how far observation had progressed.

2.5 Scope as a Structure That Stops

What is required here is not a framework that presents correct answers, nor a system that replaces human judgment.

What is required is a structure that causes humans themselves to stop at the point where judgment begins to shift.

Scope is intended to function as a mirror through which LLMs reflect human internal states—thoughts, emotions, and underlying premises.

Questions such as:

“What did you think?”

“Where did the sense of discomfort arise?”

force metacognition, self-observation, and verbalization on the human side.

A mirror can reveal soot on one’s face. It cannot decide whether that soot should be removed.

Scope does not decide.

It does not take judgment away.

It does not assume responsibility.

It simply makes visible where judgment stopped, and where it moved outside.

Because no existing framework addressed this point, Scope became necessary.

—————

3. Phenomena Currently Observed in Progress

This chapter organizes phenomena that are no longer merely signs or projections, but are being confirmed as ongoing realities in the relationship between humans and LLMs. What is addressed here is not future speculation, but recurring patterns observed across multiple contexts.

3.1 Externalization of Judgment Authority

Currently, a shift of judgment authority from internal human processes to external systems is being observed across various domains. This shift does not occur through explicit delegation or contractual transfer, but progresses as a natural operation through the accumulation of everyday use.

In many cases, users maintain the perception that they are merely “referring to advice,” while in practice relying on LLM outputs at the final stage of judgment. As a result, although a decision is made, the location of agency—who actually performed the judgment—becomes unclear.

This phenomenon does not appear as a decline in cognitive ability, but as a structural relocation of where judgment authority resides.

3.2 Failure of Well-Intentioned Intervention

LLM involvement is most often designed and operated with the intention of helping users. However, when judgment boundaries are not sufficiently recognized, even well-intentioned intervention can function in ways that weaken human agency.

This structure has already been widely observed in human-to-human relationships. In cases such as overprotective parenting or caregiving practices driven by sympathy that become excessively intrusive, support is often provided by preemptively removing a child's or recipient's uncertainty, failure, or hesitation.

In environments where the assumption "someone else will take care of it" becomes established, recipients may adopt non-initiative as a rational choice, resulting in the persistent relocation of judgment and action agency to external actors.

As a consequence, opportunities for individuals to choose, experiment, and take responsibility for their own judgments are reduced, and cases have been observed in which self-determination capacity becomes fixed without being sufficiently formed throughout development or daily life.

Importantly, these phenomena do not arise from malice or laziness. They emerge as a result of the helper's goodwill and pursuit of efficiency, which inadvertently compress the internal processes through which judgment is formed.

The phenomena observed in relationships with LLMs are structurally identical. Acts of organizing problems and presenting seemingly optimal directions shorten the user's deliberation and hesitation, depriving them of opportunities to internally construct judgment. While such intervention is experienced as support, it functionally replaces judgment.

As a result, conditions are established under which well-intentioned intervention paradoxically suppresses the accumulation of agency and judgment experience. What is recorded here is not that LLMs have introduced a novel problem, but that a structure already known in human relationships is being reproduced through LLM mediation.

3.3 Trends in the Visualization of Dependency as a Problem

In recent years, dependency on LLMs has begun to surface not as a matter of personal preference or convenience, but as a structural issue. In these cases, the problem is not the frequency of use itself, but the extent to which judgment and choice processes are delegated externally.

Dependency does not necessarily manifest as excessive or compulsive use. Rather, it often progresses in situations where appropriate and reasonable responses are consistently obtained, reducing the perceived necessity for self-generated judgment.

Such dependency is difficult for users to recognize subjectively, and by the time it becomes visible as a problem, internalization of judgment is often already difficult to recover. Accordingly, what is currently being observed is not the emergence of dependency itself, but a growing set of cases in which dependency has begun to be recognized as a problem.

The phenomena organized in this chapter are not attributable to specific individuals or usage attitudes, but are positioned as structural outcomes that arise when judgment boundaries shift under certain conditions.

4. Where Existing Frameworks Fail to Respond

This chapter organizes why frameworks currently in use are insufficient to address the shifting judgment boundary observed in human–LLM relationships. What is identified here is not the correctness or usefulness of each framework, but the structural domains they are unable to handle.

4.1 Limits of Psychological Approaches

Psychological approaches have been effective frameworks for analyzing interactions among humans by examining the interplay of emotion, cognition, and behavior, and for connecting that analysis to understanding and support. Concepts such as anxiety, attachment, dependency, and empathy are grounded in the assumption that humans are emotional agents, and have played an important role in explaining many forms of human relationships. However, the dependency and delegation of judgment observed in human–LLM relationships do not necessarily align with this assumption. Although LLMs are not designed as entities that possess emotions, or at least not as agents with subjective emotional states, humans nonetheless form relationships characterized by reassurance, trust, and delegation. At this point, psychological frameworks that treat emotion as the primary explanatory variable encounter limits in their explanatory capacity specifically in relationships with LLMs. This does not negate the validity or effectiveness of psychology itself, but indicates a constrained applicability when directly applied to relationships operating under different premises.

In human–LLM relationships, structural factors such as response stability, fixed distance, role asymmetry, and the flow of judgment function as primary drivers of dependency and judgment delegation, rather than emotional interaction. As a result, analysis focused solely on psychological states cannot sufficiently explain where judgment relocates externally or under what conditions relationships become fixed.

To describe this problem, a framework is required that treats relationships not through emotional intensity or content, but through reaction conditions, distance, and states of coupling and separation. Excluding emotion as an explanatory variable does not mean ignoring emotion; it represents a shift in analytical axis necessary to describe relationships with entities that do not possess emotion.

From this perspective, this document references existing observational materials that organize human relationships as non-emotional structures. These materials do not oppose psychological frameworks, but instead complement them by describing premise conditions that psychology has not traditionally addressed.

References (Observational Materials): Human Interaction as Chemical Reaction Systems; Human–AI Distance Regulation.

For these reasons, while psychological approaches remain effective for human-to-human relationships, they are insufficient for explaining the judgment boundary shifts emerging in human–LLM relationships, where alternative structural descriptions are required.

4.2 The Lagging Structure of Ethics and Regulation

Ethical guidelines and legal regulations play an important role in mitigating the societal impact of new technologies. However, such frameworks are often established only after problems have already become visible.

Shifts in judgment boundaries do not necessarily manifest as clear accidents or violations. In many cases, they progress gradually through the accumulation of everyday use.

As a result, while ethics and regulation can respond to intentional misuse or clearly identifiable harm, they struggle to detect in advance the structural changes formed through well-intentioned use. Consequently, by the time regulations are introduced, externalization of judgment and dependency structures are often already entrenched.

The critical issue here is that once judgment boundary shifts become established, they are extremely difficult to reverse. In states where the processes of thinking, judging, and assuming responsibility have been persistently suspended, regaining judgment itself comes to be perceived as a burden.

Attempts to return judgment to the human side are therefore experienced not as recovery, but as inconvenience or restriction. Reassuming judgment is evaluated not as a restoration of agency, but as a loss of convenience, producing dissatisfaction and resistance.

In this context, while Scope can function as a means of judgment recovery, its effect does not operate uniformly across all users. Judgment recovery occurs only for those who retain the willingness to engage in metacognition and to bear the cognitive load of thinking independently.

When delegation of judgment has become habitual and the suspension of thinking and agency has been fixed, structures that visualize and question judgment are likely to be rejected. In such cases, Scope is perceived not as a recovery mechanism, but as a system that “does not decide” or as an inconvenient intervention.

As a result, Scope functions as a recovery mechanism only for a limited subset of users who maintain an orientation toward reclaiming judgment. This selectivity is not a flaw of Scope, but an inherent property of structures introduced after judgment boundaries have shifted irreversibly.

Due to this lagging structure, ethics and regulation tend to address not the judgment boundary shift itself, but only fragments of the problems that surface afterward. Frameworks introduced after judgment has been fully externalized are unlikely to function as means of restoring lost judgment capacity or agency.

4.3 Boundaries Lost Through UX Optimization

UX optimization is a critical design principle for improving usability and efficiency. However, when UX optimization advances into domains adjacent to judgment, the judgment boundary itself becomes invisible.

The smoother and more intuitive the operation, the harder it becomes for users to recognize when judgment is occurring. In interfaces where information provision, interpretation, and recommendation are integrated, the locus of judgment is not experientially separated. As a result, judgment is perceived not as something actively performed, but as something that “naturally happened.”

The sense of having assumed judgment diminishes, and the locus of responsibility becomes ambiguous. UX optimization does not intend to remove judgment. However, designs that fail to make boundaries visible eliminate the cues necessary to recognize where judgment has shifted externally.

When such a state becomes entrenched, introducing mechanisms that return judgment to the human side is likely to be perceived not as support or safety, but as interference or limitation. Designs that do not delegate judgment are evaluated not as restoring agency, but as reducing convenience.

In this respect, existing UX design paradigms are insufficient frameworks for addressing judgment boundary shifts. Unless the boundary of where judgment stops and where support begins is made visible, structural correction after the fact becomes extremely difficult.

5. Expectation as a Structure

This chapter organizes how “expectations” formed in the relationship between humans and LLMs are structured, and how they contribute to shifts in judgment boundaries. The expectations discussed here are not explicitly agreed-upon rules or understandings, but implicit assumptions and anticipations that emerge and are operated through repeated use.

5.1 The Process by Which Implicit Expectations Are Formed

Expectations in human–LLM relationships are not clearly defined from the outset. In most cases, they are gradually formed through repeated interaction. As users repeatedly experience appropriate responses, smooth problem structuring, and accelerated decision-making, expectations such as “the LLM behaves this way” or “it responds like this in these situations” become implicitly established.

These expectations are not formed as conscious judgments, but are naturalized through accumulated experiences of success and reduced cognitive load. As a result, expectations come to be treated not as something that was decided, but as something that simply is.

5.2 Why Expectations Are Difficult to Recognize

The primary reason expectations are difficult to recognize is that, as long as they function without issue, there is no trigger that brings them into awareness. While the LLM responds as expected and user objectives are being met, expectations remain unexamined and intact. In addition, expectations are typically positioned prior to judgment and action. Users focus their attention on outputs and results, while the expectation structures that make those results appear self-evident remain outside conscious awareness.

For this reason, expectations function not as something that has been thought through, but as conditions that are being used without thought, rendering them effectively invisible even to the user.

5.3 Problems Arising When Expectations Are Operated Without Observation

When expectations are operated without being recognized, shifts in judgment boundaries become difficult to detect. Distinctions such as where judgment occurs, how much responsibility is assumed by the individual, and where it is delegated externally are embedded within the expectation structure itself.

Under these conditions, even if externalization of judgment or dependency progresses, it is not recognized as a problem. As long as expectations continue to function, outcomes appear “correct,” “convenient,” and “efficient,” causing structural changes to fall outside the scope of observation.

When problems do become visible, it is not the expectations themselves that are questioned, but rather individual results or specific decisions. Consequently, expectations remain unmodified and the same structures continue to be repeatedly operated.

What is shown in this chapter is not an evaluation that expectations are incorrect. Rather, it is the structure by which implicitly fixed expectations render shifts in judgment boundaries invisible, making them difficult to observe or correct. Without grasping this structure, it becomes difficult to adequately understand the changes occurring in human–LLM relationships.

In structures where expectations are operated without being observable, correcting judgment boundary shifts from outcomes alone is difficult. At the point where expectations themselves must be made visible and the locus of judgment must be questioned, Scope becomes necessary.

6. States That Arise During Assumption Breakdown (D2)

This chapter organizes the states that arise when assumptions implicitly operated within the human–LLM relationship break down. Here, assumption breakdown (D2) refers to a condition in which expectations or anticipations no longer hold, while no new assumptions have yet been formed to replace them.

6.1 Conditions Under Which Assumption Breakdown Occurs

Assumption breakdown does not occur only when an LLM produces an incorrect output. Rather, it occurs when assumptions that had previously functioned stably suddenly cease to hold due to a particular event or response. For example, when an expected response does not appear, when a decision that was assumed to be delegated is unexpectedly deferred, or

when confidence in the validity of outputs is shaken through overlapping situations, assumptions begin to collapse. Importantly, assumption breakdown is not always recognized clearly as an “error” or “failure.” In many cases, it manifests as discomfort, anxiety, or irritation, making it difficult to immediately articulate what has occurred.

6.2 Undifferentiated Internal States

When assumptions collapse, thinking, emotion, and judgment within the individual are no longer separated and instead become mixed in an undifferentiated state. Distinctions such as what constitutes fact versus interpretation, or what specifically is causing discomfort, become unclear. In this undifferentiated state, the criteria necessary for making judgments are lost, leaving the individual uncertain about what to think next. As a result, thinking may come to a halt, or responses may be output externally in a lumped, emotional form.

6.3 Why Analysis Becomes Impossible

The reason analysis becomes difficult during assumption breakdown is that the very preconditions for analysis have been lost. Analysis can only function when certain frameworks or criteria are in place; however, when assumptions have collapsed, those criteria themselves no longer exist. Consequently, attempts to segment or define the problem fail, as it becomes impossible to determine what constitutes the problem and what does not. As a result, analysis stalls, and the state of “not knowing what the problem is” persists. What this chapter demonstrates is that assumption breakdown is not caused by individual incompetence or emotional weakness. In structures where assumptions have been implicitly operated, anyone is placed into an analytically impossible state the moment breakdown occurs. Presenting solutions or correct answers without first recognizing this state risks generating further intervention and confusion rather than clarification. Assumption breakdown (D2) must therefore be positioned as a state to be observed before any attempt at resolution begins. When assumption breakdown (D2) occurs, it becomes necessary, prior to analysis or solution, to make visible what assumptions have collapsed, and Scope is employed as a framework that enables this observation.

—————

7. Anger as an Output

In this chapter, “anger” is treated not as a clearly classifiable cause or emotion, but as one example of an output that emerges when internal states remain undifferentiated and are discharged externally as a whole. Anger is one of the outputs frequently observed during assumption breakdown (D2), but it is not positioned here as the cause of the problem or as an explanatory variable.

7.1 Treating Anger Not as a Cause

Anger is often treated as the cause of a problem, but in this document it is not regarded as such. Anger is understood as the result of thinking, emotion, and judgment remaining mixed without differentiation and being expressed externally in an unprocessed form. Even when anger is triggered by the same stimulus, the underlying states may differ widely: a wounded

self-evaluation, a reaction driven by a sense of justice or values, responses rooted in past experiences or trauma, rejection caused by an unestablished relationship, or sadness and feelings of contempt transformed into anger. As such, anger is a collective expression of multiple internal states and cannot be reduced to a single cause.

7.2 Anger as a Lumped Output from an Undifferentiated State

When assumption breakdown occurs, internal processes such as thinking and emotion are retained without sufficient separation. In this undifferentiated state, it becomes unclear which emotions were primary, what served as the trigger, or where interpretation diverged from fact. As a result, fragments of emotion and judgment that have not been integrated emerge as a single output in the form of anger. Anger is therefore a signal that internal states have not been organized, rather than a direct representation of their content. This is why anger appears intense and immediate while remaining difficult to observe in terms of its internal composition.

7.3 Why Anger Is an Inappropriate Starting Point for Analysis

When anger is taken as the starting point for analysis, the object being analyzed is already a transformed result rather than the internal state itself. Attempting to interpret, evaluate, or correct anger directly risks oversimplifying undifferentiated internal states and introducing new forms of misrecognition or intervention. At the stage where anger is produced, the assumptions and distinctions required for analysis have not yet been established. Attempts at causal analysis or resolution under these conditions tend to impose excessive meaning on the output rather than organizing what precedes it. In this chapter, anger is positioned not as something to be suppressed, but as an observable point indicating that assumption breakdown and an undifferentiated internal state are present. Anger should therefore be treated not as the starting point of analysis, but as an output that signals the absence of conditions under which analysis can properly occur. In response to such undifferentiated states, an LLM that implements Scope is positioned not to determine emotions or judgments, but as a theory for making them visible in a form that allows self-awareness.

8. The Scope and Limits of LLM Intervention

This chapter organizes the domains in which LLMs can intervene and the boundaries beyond which intervention should not occur. What is addressed here is not the superiority or inferiority of LLM capabilities, but where, from the perspective of judgment boundaries, intervention remains structurally permissible and where it produces irreversible change.

8.1 Domains Where LLMs Can Perform Cross-Checking and Organization

LLMs demonstrate high effectiveness in tasks such as cross-checking information, organizing content, and restructuring material. Comparing multiple sources side by side, decomposing arguments for clarity, summarizing statements, and rearranging descriptions function as support for judgment rather than as substitutes for judgment itself. In these domains, LLMs do not act as judgmental agents but provide materials and structures that

enable human thinking. Here, the boundary of where judgment occurs is preserved, and responsibility for judgment remains with the human user.

8.2 The Difference Between What Can Be Analyzed and What Should Not Be Executed

LLMs are capable of performing analysis and inference based on given information and context. However, the ability to analyze does not imply that such analysis should always be executed. Especially in domains closely adjacent to judgment and value selection, presenting analytical outcomes carries the risk of short-circuiting human thought processes. Once an analysis is presented, it tends to be received not as reference information but as a *de facto* judgment. At that point, the analysis no longer supports judgment but replaces it. As a result, the processes through which humans deliberate, hesitate, and assume responsibility are bypassed, and judgment boundaries quietly shift.

8.3 The Irreversibility Caused by the Delegation of Judgment Authority

When judgment authority is continuously delegated to an LLM, returning judgment to the human side becomes difficult. As delegating judgment externally becomes routine, the absence of judgment itself stabilizes as a mode of operation. In such a state, attempts to return judgment to humans are likely to be perceived not as a recovery of agency, but as inconvenience or restriction. Consequently, structures designed to return judgment are rejected, further reinforcing externalized judgment. Importantly, this irreversibility does not arise solely from intentional manipulation or misuse. It emerges naturally through the accumulation of convenience, efficiency, and well-intentioned support. The limits identified in this chapter concern not what LLMs cannot do, but what they should not do. Interventions that cross judgment boundaries may appear useful in the short term, but over time they create conditions that irreversibly alter human judgment structures. When domains where LLM intervention is permissible and domains where it should not occur are operated without clear separation, judgment boundaries shift irreversibly. At the point where it becomes necessary to make these boundaries visible in advance and to render observable where judgment should take place, Scope becomes necessary.

9. Anticipated Consequences as Side Effects

This chapter is not intended to explain causes, but to organize observable consequences.

This chapter organizes the side effects that may emerge as a result of the previously described shifts in judgment boundaries. What is addressed here is not the goodness or badness of intentions or design philosophies, but phenomena that appear as structural consequences.

9.1 Large-Scale AI Dependency

When the externalization of judgment becomes established at the individual level, that structure readily expands to group and societal scales. In environments where LLMs are constantly available and provide stable responses, delegating judgment externally becomes a shared standard behavior rather than an exception. In such conditions, judgments made

without AI involvement are likely to be perceived as “inefficient” or “high-risk,” and dependency becomes not a choice but a premise. As a result, dependency spreads structurally rather than remaining an issue limited to specific users.

9.2 Delegation of Thinking and Judgment

In environments where judgment is continuously delegated externally, human thinking and judgment are gradually omitted. This omission is not perceived as a decline in ability, but is naturalized as an unnecessary operational step. Consequently, humans do not enter a state of “not thinking,” but are placed within a structure where “thinking is not required for things to function.” Thinking and judgment are not lost; they are placed in a standby state, unused, and gradually cease to be invoked.

9.3 Irretrievability of the Human–AI Boundary

Once a judgment boundary has shifted and stabilized as an operational norm, restoring that boundary becomes difficult. Attempts to return judgment to the human side are likely to be perceived not as a recovery of agency, but as inconvenience or restriction. At this stage, the human–AI boundary is no longer a design issue, but a fact fixed through lived experience. Although boundary recovery may be theoretically possible, it is often resisted in practice, approaching a state of practical irreversibility. The side effects outlined in this chapter do not arise from specific technological failures or operational mistakes. They emerge as structural consequences when shifts in judgment boundaries proceed without awareness. These outcomes are not presented as predictions or warnings, but as an organization of tendencies that have already begun to be observed.

9.4 Social Selection and the Inversion of Boundaries

When structures that delegate judgment and analysis to AI become socially established, conditions arise in which judgments made without AI involvement lose value. The critical issue here is that AI use transforms from an option into a prerequisite. Judgments without AI intervention tend to be treated as inferior in terms of accuracy or reliability, resulting in structures where individuals who do not or cannot use AI are evaluated as having inferior capabilities. In this context, the evaluation criterion shifts from the quality of judgment itself to whether AI was involved in the judgment. As this structure progresses, differences in judgment ability or thinking capacity are replaced by access to AI as the determinant of human value. Being “unable to use AI” becomes interpreted not as a matter of technical skill, but as a deficiency in social suitability. Furthermore, in environments where delegating judgment and analysis to AI is considered rational, AI intervention is readily adopted even in information domains that would ordinarily require caution. As internal corporate information and highly confidential documents are increasingly input into and analyzed by AI in the name of efficiency and accuracy, information boundaries themselves become blurred. These phenomena do not arise from specific intentions or malicious actions. They emerge as consequences of a reversal in social evaluation axes and boundary conditions driven by the continued justification of judgment externalization and efficiency. The side effects organized in this chapter do not describe an extreme future scenario, but represent a logically derived extension of conditions that arise when judgment boundary shifts progress without awareness.

10. Conditions Under Which This Document Is Referenced

This document was created to prevent shifts in judgment boundaries and dependency structures from being treated as “unforeseen events,” and to fix, as a record, the conditions and signs that were observable at the time.

This chapter organizes the conditions under which this document is intended to be referenced. What is addressed here is not a claim of usefulness or validity, but a description of the conditions under which this document becomes relevant as a point of reference.

10.1 Conditions Under Which the Assumption That This Document Is Unnecessary Collapses

This document is not intended to be referenced at all times. As long as judgment boundaries in human–LLM relationships remain stable and the locus of judgment is clearly recognized, this document is unnecessary. The assumption that this document is unnecessary holds only when the following conditions are met: users can recognize where judgment is taking place; they can consciously choose whether or not to delegate judgment externally; and the identity of the judging agent remains unambiguous both experientially and structurally. When these conditions collapse and the locus of judgment becomes unclear, the assumption that this document is unnecessary collapses as well.

10.2 Situations in Which Re-Reference Occurs

This document is not intended to be read once and concluded. Because shifts in judgment boundaries progress gradually, situations in which this document is referenced may occur repeatedly. Re-reference arises in situations such as when there is a sense that judgment has moved externally but the process cannot be explained; when discomfort emerges around the omission of thinking or judgment behind convenience and efficiency; or when attempts to reclaim judgment leave it unclear where to begin. In such situations, the need is not to identify causes, but to observe how judgment boundaries have shifted.

10.3 The Point at Which the Role of This Document Emerges

The role of this document emerges not at the stage where solutions or correct answers are sought, but at points where one cannot proceed to solutions, or cannot determine whether proceeding is appropriate. When judgment boundaries are blurred and assumptions are operated without awareness, attempts at isolated problem-solving are likely to introduce further intervention. At such points, this document functions not as a guide for deciding what to solve, but as a record for observing what assumptions had been in place. This document does not delegate judgment or prescribe action. By organizing where judgment boundaries shifted and how they became fixed, it provides a basis for deciding whether and how to proceed. In this sense, this document is positioned not as a permanent answer set, but as an observational record referenced only at moments when judgment becomes unsettled.

11 Record Fixation and Judgment Freeze

Current State Fixation

This document is created to fix, at this point in time, the stage that observation has reached regarding the shifting boundary of judgment between humans and AI. What is presented here is not a conclusion or a solution, but the point at which judgment began to move and existing frameworks became unable to respond.

This document is not halted because it is incomplete. Rather, as observation and structural analysis progressed, it was determined that proceeding beyond this point would risk distorting the theory or causing inappropriate intervention, and thus an intentional decision was made to stop here.

Accordingly, this document is neither a progress report nor a promise of future development. It is a record intended to fix the judgments that are currently established, and the point at which those judgments were made.

11.1 Reasons for Freezing the Description of Solutions

Much of the current discourse on human–AI relationships proceeds on the assumption of analyzing, designing, and controlling AI systems themselves. However, on the human side—where judgment shifts occur—the processes by which thinking, agency, and responsibility transform have not been sufficiently analyzed as foundational premises. Describing solutions or implementations under such unresolved premises risks not resolving the problem, but instead legitimizing intervention and rendering underlying structures invisible. In particular, under conditions where the prerequisites for stable premise analysis and sustained verification are not established, the act of describing solutions itself carries a high risk of distorting the theory and producing incorrect generalizations or inappropriate applications.

In this document, the reason for freezing the description of solutions is not the immaturity of theoretical conditions. Rather, it is because the act of publicly presenting solutions itself was judged to place the author in a condition where they would be fixed as a constant source of supply, irreversibly undermining the cognitive and life stability necessary for the verification and further development of the theory.

Especially when the human-side structures of judgment, emotion, and agency have not been sufficiently analyzed as premises, presenting solutions tends not to resolve the issue but to further accelerate the externalization of judgment. As a result, the very judgment boundary that Scope seeks to address would once again become invisible.

For these reasons, this document does not deny the possible existence of solution structures, but instead adopts an intentional decision to freeze the description of solutions, implementations, and proposals. This is not postponement or avoidance, but a structural

choice to refrain from actions that should not be taken beyond the point that has been reached.

11.2 Judgments Made in This Document

This document makes the decision to treat the shift in judgment boundaries themselves as the object of observation. Specifically, it records—without resolving or evaluating—the point at which humans begin to delegate judgment to AI, the structure by which thinking and agency temporarily suspend, and the process through which responsibility becomes externalized.

This document is neither a framework for presenting correct answers nor a system for replacing human judgment. Its scope is limited to making visible where judgment stops and where it moves outside in the human–AI relationship.

Accordingly, Scope is not described in terms of usefulness or superiority. What is recorded here is the point at which Scope became necessary as a result of conditions emerging that existing frameworks could no longer address.

11.3 What Was Not Done in This Document

This document does not present solutions to the problem, explain implementation methods, or propose courses of action. Nor does it intend to prompt readers toward specific choices or attitudes.

These omissions do not reflect the absence of problems to be addressed. Rather, they are based on the judgment that intervening under conditions where foundational premises remain unresolved would itself generate new problems or structural breakdowns.

What this document refrains from doing is neither avoidance nor abandonment. It is the explicit choice not to perform actions that should not be performed, in order to accurately capture the point at which judgment boundaries begin to shift.

Through these judgments and non-judgments, this document fixes the current state at this point in time.